

Mask R-CNN - Inspect Trained Model

Code and visualizations to test, debug, and evaluate the Mask R-CNN model.

```
In [26]: import os
import sys
import random
import math
import re
import time
import numpy as np
import tensorflow as tf
import matplotlib
import matplotlib.pyplot as plt
import matplotlib.patches as patches
import skimage

# Root directory of the project
ROOT_DIR = os.path.abspath("../")

# Import Mask RCNN
sys.path.append(ROOT_DIR) # To find local version of the library
from mrcnn import utils
from mrcnn import visualize
from mrcnn.visualize import display_images
import mrcnn.model as modellib
from mrcnn.model import log
from mrcnn.config import Config
from mrcnn import utils
import mrcnn.model as modellib
from mrcnn import visualize
from mrcnn.model import log

%matplotlib inline

# Directory to save logs and trained model
MODEL_DIR = os.path.join(ROOT_DIR, "logs")

# Local path to trained weights file
COCO_MODEL_PATH = os.path.join(ROOT_DIR, "mask_rcnn_coco.h5")
# Download COCO trained weights from Releases if needed
if not os.path.exists(COCO_MODEL_PATH):
    utils.download_trained_weights(COCO_MODEL_PATH)

# Path to Shapes trained weights
SHAPES_MODEL_PATH = os.path.join(ROOT_DIR, "logs/ped20181117T1835/mask_rcnn_ped_0001.h5")
```

Configurations

```
In [5]: class PedConfig(Config):
        """Configuration for training on the toy shapes dataset.
        Derives from the base Config class and overrides values specific
        to the toy shapes dataset.
        """
        # Give the configuration a recognizable name
        NAME = "ped"

        # Train on 1 GPU and 8 images per GPU. We can put multiple images
        on each
        # GPU because the images are small. Batch size is 8 (GPUs * image
        s/GPU).
        GPU_COUNT = 1
        IMAGES_PER_GPU = 2

        # Number of classes (including background)
        NUM_CLASSES = 1 + 3 # background + 3 shapes

        # Use small images for faster training. Set the limits of the sma
        ll side
        # the large side, and that determines the image shape.
        IMAGE_MIN_DIM = 480
        IMAGE_MAX_DIM = 640
        IMAGE_CHANNEL_COUNT = 3
        USE_MINI_MASK = False

        # Use smaller anchors because our image and objects are small
        RPN_ANCHOR_SCALES = (8, 16, 32, 64, 128) # anchor side in pixels

        # Reduce training ROIs per image because the images are small and
        have
        # few objects. Aim to allow ROI sampling to pick 33% positive ROI
        s.
        TRAIN_ROIS_PER_IMAGE = 32

        # Use a small epoch since the data is simple
        STEPS_PER_EPOCH = 100

        # use small validation steps since the epoch is small
        VALIDATION_STEPS = 5

        config = PedConfig()
        config.display()
```

Configurations:

BACKBONE	resnet101
BACKBONE_STRIDES	[4, 8, 16, 32, 64]
BATCH_SIZE	2
BBOX_STD_DEV	[0.1 0.1 0.2 0.2]
COMPUTE_BACKBONE_SHAPE	None
DETECTION_MAX_INSTANCES	100
DETECTION_MIN_CONFIDENCE	0.7
DETECTION_NMS_THRESHOLD	0.3
FPN_CLASSIF_FC_LAYERS_SIZE	1024
GPU_COUNT	1
GRADIENT_CLIP_NORM	5.0
IMAGES_PER_GPU	2
IMAGE_CHANNEL_COUNT	3
IMAGE_MAX_DIM	640
IMAGE_META_SIZE	16
IMAGE_MIN_DIM	480
IMAGE_MIN_SCALE	0
IMAGE_RESIZE_MODE	square
IMAGE_SHAPE	[640 640 3]
LEARNING_MOMENTUM	0.9
LEARNING_RATE	0.001
LOSS_WEIGHTS	{'rpn_class_loss': 1.0, 'rpn_bbox_loss': 1.0, 'mrcnn_class_loss': 1.0, 'mrcnn_bbox_loss': 1.0, 'mrcnn_mask_loss': 1.0}
MASK_POOL_SIZE	14
MASK_SHAPE	[28, 28]
MAX_GT_INSTANCES	100
MEAN_PIXEL	[123.7 116.8 103.9]
MINI_MASK_SHAPE	(56, 56)
NAME	ped
NUM_CLASSES	4
POOL_SIZE	7
POST_NMS_ROIS_INFERENCE	1000
POST_NMS_ROIS_TRAINING	2000
PRE_NMS_LIMIT	6000
ROI_POSITIVE_RATIO	0.33
RPN_ANCHOR_RATIOS	[0.5, 1, 2]
RPN_ANCHOR_SCALES	(8, 16, 32, 64, 128)
RPN_ANCHOR_STRIDE	1
RPN_BBOX_STD_DEV	[0.1 0.1 0.2 0.2]
RPN_NMS_THRESHOLD	0.7
RPN_TRAIN_ANCHORS_PER_IMAGE	256
STEPS_PER_EPOCH	100
TOP_DOWN_PYRAMID_SIZE	256
TRAIN_BN	False
TRAIN_ROIS_PER_IMAGE	32
USE_MINI_MASK	False
USE_RPN_ROIS	True
VALIDATION_STEPS	5
WEIGHT_DECAY	0.0001

```

In [13]: class PedDataset(utils.Dataset):
    """Generates the shapes synthetic dataset. The dataset consists o
    f simple
        shapes (triangles, squares, circles) placed randomly on a blank s
        urface.
        The images are generated on the fly. No file access required.
        """

    def load_ped(self, load_path="../Datasets/", istrain=True, st_ind
ex=None, batch_size=20):
    """Generate the requested number of synthetic images.
    count: number of images to generate.
    height, width: the size of the generated images.
    """

    # Add classes
    self.add_class("ped", 1, "bike")
    self.add_class("ped", 2, "car")
    self.add_class("ped", 3, "person")

    if istrain is True:
        self.X = np.load(load_path+'X_train.npy')
        self.X_mask = np.load(load_path+'X_mask_train.npy')
        self.y = np.load(load_path+'y_train.npy')

    if istrain is False:
        self.X = np.load(load_path+'X_val.npy')
        self.X_mask = np.load(load_path+'X_mask_val.npy')
        self.y = np.load(load_path+'y_val.npy')

    self.N = self.X.shape[0]

    if istrain:
        self.X = self.X[st_index:st_index+batch_size]
        self.X_mask = self.X_mask[st_index:st_index+batch_size]
        self.y = self.y[st_index:st_index+batch_size]

    for i in range(self.X.shape[0]):
        self.add_image("ped", image_id=i, path=None)

    def load_image(self, image_id, istrain=True):
        """Generate an image from the specs of the given image ID.
        Typically this function loads the image from a file, but
        in this case it generates the image on the fly from the
        specs in image_info.
        """

        image = self.X[image_id]
        if image.ndim != 3:
            image = skimage.color.gray2rgb(image)
        return image

    def image_reference(self, image_id):
        """Return the shapes data of the image."""
        return self.X[image_id].shape

    def load_mask(self, image_id):

```

```

        """Generate instance masks for shapes of the given image ID.
        """
        mask = np.zeros((self.X_mask.shape[1], self.X_mask.shape[2],
1))
        mask[:, :, 0] = np.logical_not(self.X_mask[image_id]).astype(
np.bool_)
#         print('masks are bool!')
        class_ids = np.array([self.y[image_id]+1])
        return mask, class_ids.astype(np.int32)

    def prepare(self):
        """Prepares the Dataset class for use.

        TODO: class map is not supported yet. When done, it should ha
ndle mapping
            classes from different datasets to the same class ID.
        """

        self.num_classes = int(max(self.y[:]) + 2)
        self.class_ids = np.arange(self.num_classes)
        self.class_names = ["BG", "BIKE", "CAR", "PERSON"]
        self.num_images = self.X.shape[0]
        self._image_ids = np.arange(self.num_images)

        # Mapping from source class and image IDs to internal IDs
        self.class_from_source_map = {"{ }.{ }".format(info['source'],
info['id']): id
                                     for info, id in zip(self.class_
info, self.class_ids)}
        self.image_from_source_map = {"{ }.{ }".format(info['source'],
info['id']): id
                                     for info, id in zip(self.image_
info, self.image_ids)}

        # Map sources to class_ids they support
        self.sources = list(set([i['source'] for i in self.class_info
]))
        self.source_class_ids = {}
        # Loop over datasets
        for source in self.sources:
            self.source_class_ids[source] = []
            # Find classes that belong to this dataset
            for i, info in enumerate(self.class_info):
                # Include BG class in all datasets
                if i == 0 or source == info['source']:
                    self.source_class_ids[source].append(i)

```

```
In [14]: # Run one of the code blocks

config = PedConfig()

# Shapes toy dataset
# import shapes
# config = shapes.ShapesConfig()

# MS COCO Dataset
# import coco
# config = coco.CocoConfig()
# COCO_DIR = "path to COCO dataset" # TODO: enter value here
```

```
In [15]: # Override the training configurations with a few  
# changes for inferencing.  
class InferenceConfig(config.__class__):  
    # Run detection on one image at a time  
    GPU_COUNT = 1  
    IMAGES_PER_GPU = 1  
  
config = InferenceConfig()  
config.display()
```


Configurations:

BACKBONE	resnet101
BACKBONE_STRIDES	[4, 8, 16, 32, 64]
BATCH_SIZE	1
BBOX_STD_DEV	[0.1 0.1 0.2 0.2]
COMPUTE_BACKBONE_SHAPE	None
DETECTION_MAX_INSTANCES	100
DETECTION_MIN_CONFIDENCE	0.7
DETECTION_NMS_THRESHOLD	0.3
FPN_CLASSIF_FC_LAYERS_SIZE	1024
GPU_COUNT	1
GRADIENT_CLIP_NORM	5.0
IMAGES_PER_GPU	1
IMAGE_CHANNEL_COUNT	3
IMAGE_MAX_DIM	640
IMAGE_META_SIZE	16
IMAGE_MIN_DIM	480
IMAGE_MIN_SCALE	0
IMAGE_RESIZE_MODE	square
IMAGE_SHAPE	[640 640 3]
LEARNING_MOMENTUM	0.9
LEARNING_RATE	0.001
LOSS_WEIGHTS	{'rpn_class_loss': 1.0, 'rpn_bbox_loss': 1.0, 'mrcnn_class_loss': 1.0, 'mrcnn_bbox_loss': 1.0, 'mrcnn_mask_loss': 1.0}
MASK_POOL_SIZE	14
MASK_SHAPE	[28, 28]
MAX_GT_INSTANCES	100
MEAN_PIXEL	[123.7 116.8 103.9]
MINI_MASK_SHAPE	(56, 56)
NAME	ped
NUM_CLASSES	4
POOL_SIZE	7
POST_NMS_ROIS_INFERENCE	1000
POST_NMS_ROIS_TRAINING	2000
PRE_NMS_LIMIT	6000
ROI_POSITIVE_RATIO	0.33
RPN_ANCHOR_RATIOS	[0.5, 1, 2]
RPN_ANCHOR_SCALES	(8, 16, 32, 64, 128)
RPN_ANCHOR_STRIDE	1
RPN_BBOX_STD_DEV	[0.1 0.1 0.2 0.2]
RPN_NMS_THRESHOLD	0.7
RPN_TRAIN_ANCHORS_PER_IMAGE	256
STEPS_PER_EPOCH	100
TOP_DOWN_PYRAMID_SIZE	256
TRAIN_BN	False
TRAIN_ROIS_PER_IMAGE	32
USE_MINI_MASK	False
USE_RPN_ROIS	True
VALIDATION_STEPS	5
WEIGHT_DECAY	0.0001

Notebook Preferences

```
In [16]: # Device to load the neural network on.
# Useful if you're training a model on the same
# machine, in which case use CPU and leave the
# GPU for training.
DEVICE = "/cpu:0" # /cpu:0 or /gpu:0

# Inspect the model in training or inference modes
# values: 'inference' or 'training'
# TODO: code for 'training' test mode not ready yet
TEST_MODE = "inference"
```

```
In [17]: def get_ax(rows=1, cols=1, size=16):
    """Return a Matplotlib Axes array to be used in
    all visualizations in the notebook. Provide a
    central point to control graph sizes.

    Adjust the size attribute to control how big to render images
    """

    _, ax = plt.subplots(rows, cols, figsize=(size*cols, size*rows))
    return ax
```

Load Validation Dataset

```
In [27]: # Build validation dataset
dataset = PedDataset()
dataset.load_ped(istrain=False)

# Must call before using the dataset
dataset.prepare()

print("Images: {}\nClasses: {}".format(len(dataset.image_ids), dataset.class_names))

Images: 90
Classes: ['BG', 'BIKE', 'CAR', 'PERSON']
```

Load Model

```
In [24]: # Create model in inference mode
with tf.device(DEVICE):
    model = modellib.MaskRCNN(mode="inference", model_dir=MODEL_DIR,
                               config=config)

# # Set weights file path
# if config.NAME == "shapes":
#     weights_path = SHAPES_MODEL_PATH
# elif config.NAME == "coco":
#     weights_path = COCO_MODEL_PATH
# Or, uncomment to load the last model you trained
weights_path = SHAPES_MODEL_PATH

# Load weights
print("Loading weights ", weights_path)
model.load_weights(weights_path, by_name=True)
```

```
Loading weights /home/paperspace/Documents/PedNet/logs/ped20181117T1
835/mask_rcnn_ped_0001.h5
Re-starting from epoch 1
```


Run Detection

```
In [28]: image_id = random.choice(dataset.image_ids)
image, image_meta, gt_class_id, gt_bbox, gt_mask =\
    modellib.load_image_gt(dataset, config, image_id, use_mini_mask=F
alse)
info = dataset.image_info[image_id]
print("image ID: {}.{} ({}).{}".format(info["source"], info["id"], im
age_id,
                                     dataset.image_reference(image_
id)))
# Run object detection
results = model.detect([image], verbose=1)

# Display results
ax = get_ax(1)
r = results[0]
visualize.display_instances(image, r['rois'], r['masks'], r['class_id
s'],
                           dataset.class_names, r['scores'], ax=ax,
                           title="Predictions")
log("gt_class_id", gt_class_id)
log("gt_bbox", gt_bbox)
log("gt_mask", gt_mask)
```


```
image ID: ped.63 (63) (480, 640)
Processing 1 images
image                shape: (640, 640, 3)           min: 0.00000
  max: 255.00000 float64
molded_images        shape: (1, 640, 640, 3)         min: -123.70000
  max: 151.10000 float64
image metas          shape: (1, 16)                  min: 0.00000
  max: 640.00000 int64
anchors              shape: (1, 102300, 4)           min: -0.14164
  max: 1.04149 float32
gt_class_id           shape: (1, 1)                  min: 2.00000
  max: 2.00000 int32
gt_bbox              shape: (1, 4)                   min: 0.00000
  max: 433.00000 int32
gt_mask              shape: (640, 640, 1)            min: 0.00000
  max: 1.00000 bool
```

Predictions




Precision-Recall

```
In [29]: # Draw precision-recall curve
AP, precisions, recalls, overlaps = utils.compute_ap(gt_bbox, gt_class_id, gt_mask,
                                                    r['rois'], r['class_ids'],
                                                    r['scores'], r['masks'])
visualize.plot_precision_recall(AP, precisions, recalls)
```



```
In [30]: # Grid of ground truth objects and their predictions
visualize.plot_overlaps(gt_class_id, r['class_ids'], r['scores'],
                        overlaps, dataset.class_names)
```



Compute mAP @ IoU=50 on Batch of Images

```

In [31]: # Compute VOC-style Average Precision
def compute_batch_ap(image_ids):
    APs = []
    for image_id in image_ids:
        # Load image
        image, image_meta, gt_class_id, gt_bbox, gt_mask = \
            modellib.load_image_gt(dataset, config,
                                   image_id, use_mini_mask=False)

        # Run object detection
        results = model.detect([image], verbose=0)
        # Compute AP
        r = results[0]
        AP, precisions, recalls, overlaps = \
            utils.compute_ap(gt_bbox, gt_class_id, gt_mask,
                             r['rois'], r['class_ids'], r['scores'],
                             r['masks'])
        APs.append(AP)
    return APs

# Pick a set of random images
image_ids = np.random.choice(dataset.image_ids, 10)
APs = compute_batch_ap(image_ids)
print("mAP @ IoU=50: ", np.mean(APs))

mAP @ IoU=50:  1.0

```

Step by Step Prediction

Stage 1: Region Proposal Network

The Region Proposal Network (RPN) runs a lightweight binary classifier on a lot of boxes (anchors) over the image and returns object/no-object scores. Anchors with high *objectness* score (positive anchors) are passed to the stage two to be classified.

Often, even positive anchors don't cover objects fully. So the RPN also regresses a refinement (a delta in location and size) to be applied to the anchors to shift it and resize it a bit to the correct boundaries of the object.

1.a RPN Targets

The RPN targets are the training values for the RPN. To generate the targets, we start with a grid of anchors that cover the full image at different scales, and then we compute the IoU of the anchors with ground truth object. Positive anchors are those that have an IoU ≥ 0.7 with any ground truth object, and negative anchors are those that don't cover any object by more than 0.3 IoU. Anchors in between (i.e. cover an object by IoU ≥ 0.3 but < 0.7) are considered neutral and excluded from training.

To train the RPN regressor, we also compute the shift and resizing needed to make the anchor cover the ground truth object completely.



```
In [32]: # Generate RPN trainig targets
# target_rpn_match is 1 for positive anchors, -1 for negative anchors
# and 0 for neutral anchors.
target_rpn_match, target_rpn_bbox = modellib.build_rpn_targets(
    image.shape, model.anchors, gt_class_id, gt_bbox, model.config)
log("target_rpn_match", target_rpn_match)
log("target_rpn_bbox", target_rpn_bbox)

positive_anchor_ix = np.where(target_rpn_match[:] == 1)[0]
negative_anchor_ix = np.where(target_rpn_match[:] == -1)[0]
neutral_anchor_ix = np.where(target_rpn_match[:] == 0)[0]
positive_anchors = model.anchors[positive_anchor_ix]
negative_anchors = model.anchors[negative_anchor_ix]
neutral_anchors = model.anchors[neutral_anchor_ix]
log("positive_anchors", positive_anchors)
log("negative_anchors", negative_anchors)
log("neutral anchors", neutral_anchors)

# Apply refinement deltas to positive anchors
refined_anchors = utils.apply_box_deltas(
    positive_anchors,
    target_rpn_bbox[:positive_anchors.shape[0]] * model.config.RPN_BB
    OX_STD_DEV)
log("refined_anchors", refined_anchors, )
```

```
target_rpn_match      shape: (102300,)      min:   -1.00000
  max:    1.00000  int32
target_rpn_bbox        shape: (256, 4)      min:   -6.02146
  max:    6.54337  float64
positive_anchors       shape: (1, 4)        min:   37.49033
  max:   365.25483  float64
negative_anchors       shape: (255, 4)      min:   -6.62742
  max:   646.62742  float64
neutral anchors       shape: (102044, 4)    min:  -90.50967
  max:   666.50967  float64
refined_anchors        shape: (1, 4)        min:   -0.00002
  max:   432.99994  float32
```

```
In [33]: # Display positive anchors before refinement (dotted) and  
# after refinement (solid).  
visualize.draw_boxes(image, boxes=positive_anchors, refined_boxes=ref  
ined_anchors, ax=get_ax())
```



1.b RPN Predictions

Here we run the RPN graph and display its predictions.


```
In [34]: # Run RPN sub-graph
pillar = model.keras_model.get_layer("ROI").output # node to start searching from

# TF 1.4 and 1.9 introduce new versions of NMS. Search for all names to support TF 1.3~1.10
nms_node = model.ancestor(pillar, "ROI/rpn_non_max_suppression:0")
if nms_node is None:
    nms_node = model.ancestor(pillar, "ROI/rpn_non_max_suppression/NonMaxSuppressionV2:0")
if nms_node is None: #TF 1.9-1.10
    nms_node = model.ancestor(pillar, "ROI/rpn_non_max_suppression/NonMaxSuppressionV3:0")


rpn = model.run_graph([image], [
    ("rpn_class", model.keras_model.get_layer("rpn_class").output),
    ("pre_nms_anchors", model.ancestor(pillar, "ROI/pre_nms_anchors:0")),
    ("refined_anchors", model.ancestor(pillar, "ROI/refined_anchors:0")),
    ("refined_anchors_clipped", model.ancestor(pillar, "ROI/refined_anchors_clipped:0")),
    ("post_nms_anchor_ix", nms_node),
    ("proposals", model.keras_model.get_layer("ROI").output),
])
```

rpn_class	shape: (1, 102300, 2)	min: 0.00000
max: 1.00000 float32		
pre_nms_anchors	shape: (1, 6000, 4)	min: -0.14164
max: 1.04149 float32		
refined_anchors	shape: (1, 6000, 4)	min: -87.76527
max: 88.22453 float32		
refined_anchors_clipped	shape: (1, 6000, 4)	min: 0.00000
max: 1.00000 float32		
post_nms_anchor_ix	shape: (1000,)	min: 0.00000
max: 2062.00000 int32		
proposals	shape: (1, 1000, 4)	min: 0.00000
max: 1.00000 float32		


```
In [35]: # Show top anchors by score (before refinement)
limit = 100
sorted_anchor_ids = np.argsort(rpn['rpn_class'][:, :, 1].flatten())[::-1]
visualize.draw_boxes(image, boxes=model.anchors[sorted_anchor_ids[:limit]], ax=get_ax())
```



```
In [36]: # Show top anchors with refinement. Then with clipping to image boundaries
limit = 50
ax = get_ax(1, 2)
pre_nms_anchors = utils.denorm_boxes(rpn["pre_nms_anchors"][0], image
    .shape[:2])
refined_anchors = utils.denorm_boxes(rpn["refined_anchors"][0], image
    .shape[:2])
refined_anchors_clipped = utils.denorm_boxes(rpn["refined_anchors_clipped"][0], image.shape[:2])
visualize.draw_boxes(image, boxes=pre_nms_anchors[:limit],
    refined_boxes=refined_anchors[:limit], ax=ax[0])
visualize.draw_boxes(image, refined_boxes=refined_anchors_clipped[:limit], ax=ax[1])
```



```
In [37]: # Show refined anchors after non-max suppression  
limit = 50  
ixs = rpn["post_nms_anchor_ix"][:limit]  
visualize.draw_boxes(image, refined_boxes=refined_anchors_clipped[ixs]  
], ax=get_ax())
```



```
In [38]: # Show final proposals
# These are the same as the previous step (refined anchors
# after NMS) but with coordinates normalized to [0, 1] range.
limit = 50
# Convert back to image coordinates for display
h, w = config.IMAGE_SHAPE[:2]
proposals = rpn['proposals'][0, :limit] * np.array([h, w, h, w])
visualize.draw_boxes(image, refined_boxes=proposals, ax=get_ax())
```



```
In [39]: # Measure the RPN recall (percent of objects covered by anchors)
# Here we measure recall for 3 different methods:
# - All anchors
# - All refined anchors
# - Refined anchors after NMS
iou_threshold = 0.7

recall, positive_anchor_ids = utils.compute_recall(model.anchors, gt_
bbox, iou_threshold)
print("All Anchors ({:5})      Recall: {:.3f}  Positive anchors: {}".format(
    model.anchors.shape[0], recall, len(positive_anchor_ids)))

recall, positive_anchor_ids = utils.compute_recall(rpn['refined_anchors']
[0], gt_bbox, iou_threshold)
print("Refined Anchors ({:5})  Recall: {:.3f}  Positive anchors: {}".format(
    rpn['refined_anchors'].shape[1], recall, len(positive_anchor_ids
)))

recall, positive_anchor_ids = utils.compute_recall(proposals, gt_bbox
, iou_threshold)
print("Post NMS Anchors ({:5})  Recall: {:.3f}  Positive anchors: {}".format(
    proposals.shape[0], recall, len(positive_anchor_ids)))
```

All Anchors (102300) Recall: 0.000 Positive anchors: 0
Refined Anchors (6000) Recall: 0.000 Positive anchors: 0
Post NMS Anchors (50) Recall: 1.000 Positive anchors: 3

Stage 2: Proposal Classification

This stage takes the region proposals from the RPN and classifies them.

2.a Proposal Classification

Run the classifier heads on proposals to generate class probabilities and bounding box regressions.


```
In [40]: # Get input and output to classifier and mask heads.
mrcnn = model.run_graph([image], [
    ("proposals", model.keras_model.get_layer("ROI").output),
    ("probs", model.keras_model.get_layer("mrcnn_class").output),
    ("deltas", model.keras_model.get_layer("mrcnn_bbox").output),
    ("masks", model.keras_model.get_layer("mrcnn_mask").output),
    ("detections", model.keras_model.get_layer("mrcnn_detection").out
put),
])
```

proposals		shape: (1, 1000, 4)	min: 0.00000
max:	1.00000	float32	
probs		shape: (1, 1000, 4)	min: 0.00000
max:	1.00000	float32	
deltas		shape: (1, 1000, 4, 4)	min: -5.11462
max:	4.14306	float32	
masks		shape: (1, 100, 28, 28, 4)	min: 0.00022
max:	0.99996	float32	
detections		shape: (1, 100, 6)	min: 0.00000
max:	2.00000	float32	

```
In [41]: # Get detection class IDs. Trim zero padding.
det_class_ids = mrcnn['detections'][0, :, 4].astype(np.int32)
det_count = np.where(det_class_ids == 0)[0][0]
det_class_ids = det_class_ids[:det_count]
detections = mrcnn['detections'][0, :det_count]

print("{} detections: {}".format(
    det_count, np.array(dataset.class_names)[det_class_ids]))

captions = ["{} {:.3f}".format(dataset.class_names[int(c)], s) if c >
0 else ""
            for c, s in zip(detections[:, 4], detections[:, 5])]
visualize.draw_boxes(
    image,
    refined_boxes=utils.denorm_boxes(detections[:, :4], image.shape[:
2]),
    visibilities=[2] * len(detections),
    captions=captions, title="Detections",
    ax=get_ax())
```

1 detections: ['CAR']

Detections



2.c Step by Step Detection

Here we dive deeper into the process of processing the detections.

```
In [42]: # Proposals are in normalized coordinates. Scale them
# to image coordinates.
h, w = config.IMAGE_SHAPE[:2]
proposals = np.around(mrcnn["proposals"][0] * np.array([h, w, h, w]))
.astype(np.int32)

# Class ID, score, and mask per proposal
roi_class_ids = np.argmax(mrcnn["probs"][0], axis=1)
roi_scores = mrcnn["probs"][0, np.arange(roi_class_ids.shape[0]), roi_class_ids]
roi_class_names = np.array(dataset.class_names)[roi_class_ids]
roi_positive_ixs = np.where(roi_class_ids > 0)[0]

# How many ROIs vs empty rows?
print("{} Valid proposals out of {}".format(np.sum(np.any(proposals, axis=1)), proposals.shape[0]))
print("{} Positive ROIs".format(len(roi_positive_ixs)))

# Class counts
print(list(zip(*np.unique(roi_class_names, return_counts=True))))

1000 Valid proposals out of 1000
25 Positive ROIs
[('BG', 975), ('CAR', 25)]
```

```
In [43]: # Display a random sample of proposals.
# Proposals classified as background are dotted, and
# the rest show their class and confidence score.
limit = 200
ixs = np.random.randint(0, proposals.shape[0], limit)
captions = ["{} {:.3f}"].format(dataset.class_names[c], s) if c > 0 else ""
for c, s in zip(roi_class_ids[ixs], roi_scores[ixs]):
visualize.draw_boxes(image, boxes=proposals[ixs],
visibilities=np.where(roi_class_ids[ixs] > 0, 2,
1),
captions=captions, title="ROIs Before Refinement",
ax=get_ax())
```

ROIs Before Refinement



Apply Bounding Box Refinement

```

In [44]: # Class-specific bounding box shifts.
roi_bbox_specific = mrcnn["deltas"][0, np.arange(proposals.shape[0]),
roi_class_ids]
log("roi_bbox_specific", roi_bbox_specific)

# Apply bounding box transformations
# Shape: [N, (y1, x1, y2, x2)]
refined_proposals = utils.apply_box_deltas(
    proposals, roi_bbox_specific * config.BBOX_STD_DEV).astype(np.int
32)
log("refined_proposals", refined_proposals)

# Show positive proposals
# ids = np.arange(roi_boxes.shape[0]) # Display all
limit = 5
ids = np.random.randint(0, len(roi_positive_ixs), limit) # Display r
andom sample
captions = ["{} {:.3f}".format(dataset.class_names[c], s) if c > 0 el
se ""
            for c, s in zip(roi_class_ids[roi_positive_ixs][ids], roi
_scores[roi_positive_ixs][ids])]
visualize.draw_boxes(image, boxes=proposals[roi_positive_ixs][ids],
                    refined_boxes=refined_proposals[roi_positive_ixs
][ids],
                    visibilities=np.where(roi_class_ids[roi_positive
_ixs][ids] > 0, 1, 0),
                    captions=captions, title="ROIs After Refinement"
,
                    ax=get_ax())

```

```

roi_bbox_specific      shape: (1000, 4)      min:   -3.38738
max:    3.11136 float32
refined_proposals      shape: (1000, 4)      min:  -191.00000
max:   771.00000 int32

```

ROIs After Refinement



Filter Low Confidence Detections

```

In [45]: # Remove boxes classified as background
keep = np.where(roi_class_ids > 0)[0]
print("Keep {} detections:\n{}".format(keep.shape[0], keep))

```

Keep 25 detections:

```

[ 0  1  2  3  4  6  7  8  9 10 12 14 15 17 57 79 160
161
170 186 236 258 292 437 585]

```

```
In [46]: # Remove low confidence detections
keep = np.intersect1d(keep, np.where(roi_scores >= config.DETECTION_MIN_CONFIDENCE)[0])
print("Remove boxes below {} confidence. Keep {}:\n{}".format(
    config.DETECTION_MIN_CONFIDENCE, keep.shape[0], keep))
```

Remove boxes below 0.7 confidence. Keep 23:

```
[ 0  1  2  3  4  6  7  8  9 10 12 14 15 17 57 79 160
161
170 186 236 258 585]
```

Per-Class Non-Max Suppression

```
In [47]: # Apply per-class non-max suppression
pre_nms_boxes = refined_proposals[keep]
pre_nms_scores = roi_scores[keep]
pre_nms_class_ids = roi_class_ids[keep]

nms_keep = []
for class_id in np.unique(pre_nms_class_ids):
    # Pick detections of this class
    ix = np.where(pre_nms_class_ids == class_id)[0]
    # Apply NMS
    class_keep = utils.non_max_suppression(pre_nms_boxes[ix],
                                           pre_nms_scores[ix],
                                           config.DETECTION_NMS_THRESHOLD)
    # Map indices
    class_keep = keep[ix[class_keep]]
    nms_keep = np.union1d(nms_keep, class_keep)
    print("{:22}: {} -> {}".format(dataset.class_names[class_id][:20],
                                   keep[ix], class_keep))


keep = np.intersect1d(keep, nms_keep).astype(np.int32)
print("\nKept after per-class NMS: {} \n {}".format(keep.shape[0], keep))
```

```
CAR : [ 0  1  2  3  4  6  7  8  9 10 12
14 15 17 57 79 160 161
170 186 236 258 585] -> [3]
```

```
Kept after per-class NMS: 1
[3]
```

```
In [48]: # Show final detections
ixs = np.arange(len(keep)) # Display all
# ixs = np.random.randint(0, len(keep), 10) # Display random sample
captions = ["{} {:.3f}".format(dataset.class_names[c], s) if c > 0 else ""
            for c, s in zip(roi_class_ids[keep][ixs], roi_scores[keep][ixs])]
visualize.draw_boxes(
    image, boxes=proposals[keep][ixs],
    refined_boxes=refined_proposals[keep][ixs],
    visibilities=np.where(roi_class_ids[keep][ixs] > 0, 1, 0),
    captions=captions, title="Detections after NMS",
    ax=get_ax())
```

Detections after NMS



Stage 3: Generating Masks

This stage takes the detections (refined bounding boxes and class IDs) from the previous layer and runs the mask head to generate segmentation masks for every instance.

3.a Mask Targets

These are the training targets for the mask branch

```
In [49]: display_images(np.transpose(gt_mask, [2, 0, 1]), cmap="Blues")
```



3.b Predicted Masks

```
In [50]: # Get predictions of mask head
mrcnn = model.run_graph([image], [
    ("detections", model.keras_model.get_layer("mrcnn_detection").output),
    ("masks", model.keras_model.get_layer("mrcnn_mask").output),
])

# Get detection class IDs. Trim zero padding.
det_class_ids = mrcnn['detections'][0, :, 4].astype(np.int32)
det_count = np.where(det_class_ids == 0)[0][0]
det_class_ids = det_class_ids[:det_count]


print("{} detections: {}".format(
    det_count, np.array(dataset.class_names)[det_class_ids]))

detections          shape: (1, 100, 6)          min:    0.00000
  max:    2.00000  float32
masks               shape: (1, 100, 28, 28, 4)    min:    0.00022
  max:    0.99996  float32
1 detections: ['CAR']
```

```
In [51]: # Masks
det_boxes = utils.denorm_boxes(mrcnn["detections"][0, :, :4], image.shape[:2])
det_mask_specific = np.array([mrcnn["masks"][0, i, :, :, c]
                              for i, c in enumerate(det_class_ids)])
det_masks = np.array([utils.unmold_mask(m, det_boxes[i], image.shape)
                      for i, m in enumerate(det_mask_specific)])
log("det_mask_specific", det_mask_specific)
log("det_masks", det_masks)
```

```
det_mask_specific      shape: (1, 28, 28)          min:    0.00022
max:    0.99996 float32
det_masks              shape: (1, 640, 640)        min:    0.00000
max:    1.00000 bool
```

```
In [52]: display_images(det_mask_specific[:4] * 255, cmap="Blues", interpolation="none")
```



```
In [53]: display_images(det_masks[:4] * 255, cmap="Blues", interpolation="none")
```



Visualize Activations

In some cases it helps to look at the output from different layers and visualize them to catch issues and odd patterns.


```
In [54]: # Get activations of a few sample layers
activations = model.run_graph([image], [
    ("input_image",      model.keras_model.get_layer("input_image").
    .output),
    ("res4w_out",        model.keras_model.get_layer("res4w_out").o
    utput), # for resnet100
    ("rpn_bbox",         model.keras_model.get_layer("rpn_bbox").ou
    tput),
    ("roi",              model.keras_model.get_layer("ROI").output
    ),
    ],
    1)
```

```
input_image      shape: (1, 640, 640, 3)      min: -123.70000
    max: 151.10001 float32
res4w_out        shape: (1, 40, 40, 1024)      min:  0.00000
    max:  57.82650 float32
rpn_bbox         shape: (1, 102300, 4)         min: -15.51891
    max:  205.08542 float32
roi              shape: (1, 1000, 4)           min:  0.00000
    max:    1.00000 float32
```


```
In [55]: # Input image (normalized)
_ = plt.imshow(modellib.unmold_image(activations["input_image"][0],co
nfig))
```



```
In [56]: # Backbone feature map
display_images(np.transpose(activations["res4w_out"][0,:,:,:4], [2, 0
, 1]))
```



```
In [57]: # Histograms of RPN bounding box deltas
plt.figure(figsize=(12, 3))
plt.subplot(1, 4, 1)
plt.title("dy")
_ = plt.hist(activations["rpn_bbox"][0,:,0], 50)
plt.subplot(1, 4, 2)
plt.title("dx")
_ = plt.hist(activations["rpn_bbox"][0,:,1], 50)
plt.subplot(1, 4, 3)
plt.title("dw")
_ = plt.hist(activations["rpn_bbox"][0,:,2], 50)
plt.subplot(1, 4, 4)
plt.title("dh")
_ = plt.hist(activations["rpn_bbox"][0,:,3], 50)
```



```
In [58]: # Distribution of y, x coordinates of generated proposals
plt.figure(figsize=(10, 5))
plt.subplot(1, 2, 1)
plt.title("y1, x1")
plt.scatter(activations["roi"][0,:,0], activations["roi"][0,:,1])
plt.subplot(1, 2, 2)
plt.title("y2, x2")
plt.scatter(activations["roi"][0,:,2], activations["roi"][0,:,3])
plt.show()
```

