

Mask R-CNN - Train on Shapes Dataset

This notebook shows how to train Mask R-CNN on your own dataset. To keep things simple we use a synthetic dataset of shapes (squares, triangles, and circles) which enables fast training. You'd still need a GPU, though, because the network backbone is a Resnet101, which would be too slow to train on a CPU. On a GPU, you can start to get okay-ish results in a few minutes, and good results in less than an hour.

The code of the *Shapes* dataset is included below. It generates images on the fly, so it doesn't require downloading any data. And it can generate images of any size, so we pick a small image size to train faster.

```
In [1]: import os
import sys
import random
import math
import re
import time
import numpy as np
import cv2
import matplotlib
import matplotlib.pyplot as plt
from utils.raw_data import *
import skimage

# Root directory of the project
ROOT_DIR = os.path.abspath("../")

# Import Mask RCNN
sys.path.append(ROOT_DIR) # To find local version of the library
from mrcnn.config import Config
from mrcnn import utils
import mrcnn.model as modellib
from mrcnn import visualize
from mrcnn.model import log

%matplotlib inline

# Directory to save logs and trained model
MODEL_DIR = os.path.join(ROOT_DIR, "logs")

# Local path to trained weights file
COCO_MODEL_PATH = os.path.join(ROOT_DIR, "mask_rcnn_coco.h5")
# Download COCO trained weights from Releases if needed
if not os.path.exists(COCO_MODEL_PATH):
    utils.download_trained_weights(COCO_MODEL_PATH)

/home/paperspace/anaconda3/lib/python3.6/site-packages/h5py/__init__.py:36: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.
  from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

Configurations

```

In [2]: class PedConfig(Config):
        """Configuration for training on the toy shapes dataset.
        Derives from the base Config class and overrides values specific
        to the toy shapes dataset.
        """
        # Give the configuration a recognizable name
        NAME = "ped"

        # Train on 1 GPU and 8 images per GPU. We can put multiple images
        on each
        # GPU because the images are small. Batch size is 8 (GPUs * image
        s/GPU).
        GPU_COUNT = 1
        IMAGES_PER_GPU = 2

        # Number of classes (including background)
        NUM_CLASSES = 1 + 3 # background + 3 shapes

        # Use small images for faster training. Set the limits of the sma
        ll side
        # the large side, and that determines the image shape.
        IMAGE_MIN_DIM = 480
        IMAGE_MAX_DIM = 640
        IMAGE_CHANNEL_COUNT = 3
        USE_MINI_MASK = False

        # Use smaller anchors because our image and objects are small
        RPN_ANCHOR_SCALES = (8, 16, 32, 64, 128) # anchor side in pixels

        # Reduce training ROIs per image because the images are small and
        have
        # few objects. Aim to allow ROI sampling to pick 33% positive ROI
        s.
        TRAIN_ROIS_PER_IMAGE = 32

        # Use a small epoch since the data is simple
        STEPS_PER_EPOCH = 100

        # use small validation steps since the epoch is small
        VALIDATION_STEPS = 5

        config = PedConfig()
        config.display()

```

Configurations:

BACKBONE	resnet101
BACKBONE_STRIDES	[4, 8, 16, 32, 64]
BATCH_SIZE	2
BBOX_STD_DEV	[0.1 0.1 0.2 0.2]
COMPUTE_BACKBONE_SHAPE	None
DETECTION_MAX_INSTANCES	100
DETECTION_MIN_CONFIDENCE	0.7
DETECTION_NMS_THRESHOLD	0.3
FPN_CLASSIF_FC_LAYERS_SIZE	1024
GPU_COUNT	1
GRADIENT_CLIP_NORM	5.0
IMAGES_PER_GPU	2
IMAGE_CHANNEL_COUNT	3
IMAGE_MAX_DIM	640
IMAGE_META_SIZE	16
IMAGE_MIN_DIM	480
IMAGE_MIN_SCALE	0
IMAGE_RESIZE_MODE	square
IMAGE_SHAPE	[640 640 3]
LEARNING_MOMENTUM	0.9
LEARNING_RATE	0.001
LOSS_WEIGHTS	{'rpn_class_loss': 1.0, 'rpn_bbox_loss': 1.0, 'mrcnn_class_loss': 1.0, 'mrcnn_bbox_loss': 1.0, 'mrcnn_mask_loss': 1.0}
MASK_POOL_SIZE	14
MASK_SHAPE	[28, 28]
MAX_GT_INSTANCES	100
MEAN_PIXEL	[123.7 116.8 103.9]
MINI_MASK_SHAPE	(56, 56)
NAME	ped
NUM_CLASSES	4
POOL_SIZE	7
POST_NMS_ROIS_INFERENCE	1000
POST_NMS_ROIS_TRAINING	2000
PRE_NMS_LIMIT	6000
ROI_POSITIVE_RATIO	0.33
RPN_ANCHOR_RATIOS	[0.5, 1, 2]
RPN_ANCHOR_SCALES	(8, 16, 32, 64, 128)
RPN_ANCHOR_STRIDE	1
RPN_BBOX_STD_DEV	[0.1 0.1 0.2 0.2]
RPN_NMS_THRESHOLD	0.7
RPN_TRAIN_ANCHORS_PER_IMAGE	256
STEPS_PER_EPOCH	100
TOP_DOWN_PYRAMID_SIZE	256
TRAIN_BN	False
TRAIN_ROIS_PER_IMAGE	32
USE_MINI_MASK	False
USE_RPN_ROIS	True
VALIDATION_STEPS	5
WEIGHT_DECAY	0.0001

Notebook Preferences

```
In [3]: def get_ax(rows=1, cols=1, size=8):  
        """Return a Matplotlib Axes array to be used in  
        all visualizations in the notebook. Provide a  
        central point to control graph sizes.  
  
        Change the default size attribute to control the size  
        of rendered images  
        """  
        _, ax = plt.subplots(rows, cols, figsize=(size*cols, size*rows))  
        return ax
```

Dataset

Create a synthetic dataset

Extend the Dataset class and add a method to load the shapes dataset, `load_shapes()`, and override the following methods:

- `load_image()`
- `load_mask()`
- `image_reference()`

```

In [4]: class PedDataset(utils.Dataset):
    """Generates the shapes synthetic dataset. The dataset consists of simple
    shapes (triangles, squares, circles) placed randomly on a blank surface.
    The images are generated on the fly. No file access required.
    """

    def load_ped(self, load_path="../Datasets/", istrain=True, st_index=None, batch_size=20):
        """Generate the requested number of synthetic images.
        count: number of images to generate.
        height, width: the size of the generated images.
        """

        # Add classes
        self.add_class("ped", 1, "bike")
        self.add_class("ped", 2, "car")
        self.add_class("ped", 3, "person")

        if istrain is True:
            self.X = np.load(load_path+'X_train.npy')
            self.X_mask = np.load(load_path+'X_mask_train.npy')
            self.y = np.load(load_path+'y_train.npy')

        if istrain is False:
            self.X = np.load(load_path+'X_val.npy')
            self.X_mask = np.load(load_path+'X_mask_val.npy')
            self.y = np.load(load_path+'y_val.npy')

        self.N = self.X.shape[0]

        if istrain:
            self.X = self.X[st_index:st_index+batch_size]
            self.X_mask = self.X_mask[st_index:st_index+batch_size]
            self.y = self.y[st_index:st_index+batch_size]

        for i in range(self.X.shape[0]):
            self.add_image("ped", image_id=i, path=None)

    def load_image(self, image_id, istrain=True):
        """Generate an image from the specs of the given image ID.
        Typically this function loads the image from a file, but
        in this case it generates the image on the fly from the
        specs in image_info.
        """

        image = self.X[image_id]
        if image.ndim != 3:
            image = skimage.color.gray2rgb(image)
        return image

    def image_reference(self, image_id):
        """Return the shapes data of the image."""
        return self.X[image_id].shape

    def load_mask(self, image_id):

```

```

        """Generate instance masks for shapes of the given image ID.
        """
        mask = np.zeros((self.X_mask.shape[1], self.X_mask.shape[2],
1))
        mask[:, :, 0] = np.logical_not(self.X_mask[image_id]).astype(
np.bool_)
#         print('masks are bool!')
        class_ids = np.array([self.y[image_id]+1])
        return mask, class_ids.astype(np.int32)

    def prepare(self):
        """Prepares the Dataset class for use.

        TODO: class map is not supported yet. When done, it should ha
ndle mapping
            classes from different datasets to the same class ID.
        """

        self.num_classes = int(max(self.y[:]) + 2)
        self.class_ids = np.arange(self.num_classes)
        self.class_names = ["BG", "BIKE", "CAR", "PERSON"]
        self.num_images = self.X.shape[0]
        self._image_ids = np.arange(self.num_images)

        # Mapping from source class and image IDs to internal IDs
        self.class_from_source_map = {"{ }.{ }".format(info['source'],
info['id']): id
                                     for info, id in zip(self.class_
info, self.class_ids)}
        self.image_from_source_map = {"{ }.{ }".format(info['source'],
info['id']): id
                                     for info, id in zip(self.image_
info, self.image_ids)}

        # Map sources to class_ids they support
        self.sources = list(set([i['source'] for i in self.class_info
1]))
        self.source_class_ids = {}
        # Loop over datasets
        for source in self.sources:
            self.source_class_ids[source] = []
            # Find classes that belong to this dataset
            for i, info in enumerate(self.class_info):
                # Include BG class in all datasets
                if i == 0 or source == info['source']:
                    self.source_class_ids[source].append(i)

```

```

In [5]: # Training dataset
dataset_train = PedDataset()
dataset_train.load_ped(istrain=True, st_index=0)
dataset_train.prepare()

```

```
In [6]: # Validation dataset
dataset_val = PedDataset()
dataset_val.load_ped(istrain=False)
dataset_val.prepare()
```

Ceate Model

```
In [7]: # Create model in training mode
model = modellib.MaskRCNN(mode="training", config=config,
                           model_dir=MODEL_DIR)
```


In [10]: *# Train Network*

```
num_epoch = 10
batch_size = 20
N = dataset_train.N / batch_size
lr = config.LEARNING_RATE
for e in range(num_epoch):
    for b in range(int(N)):
        dataset_train.load_ped(istrain=True, st_index=b*batch_size)
        dataset_train.prepare()
        print("Epoch %d: Training batch %d out of %d" % (e, b, int(N)))
    ))
    # Which weights to start with?
    if b == 0:
        model.load_weights(COCO_MODEL_PATH, by_name=True,
                           exclude=["mrcnn_class_logits", "mrcnn_bbox_fc",
                                   "mrcnn_bbox", "mrcnn_mask"])
    else:
        # Load the last model you trained and continue training
        model.load_weights(model.find_last(), by_name=True)

    model.train(dataset_train, dataset_val,
                 learning_rate=lr,
                 epochs=1,
                 layers='heads')

    lr = lr * 0.90
```

Epoch 0: Training batch 0 out of 40

WARNING:tensorflow:From /home/paperspace/anaconda3/lib/python3.6/site-packages/tensorflow/contrib/learn/python/learn/datasets/base.py:198: retry (from tensorflow.contrib.learn.python.learn.datasets.base) is deprecated and will be removed in a future version.

Instructions for updating:

Use the retry module or similar alternatives.

Starting at epoch 0. LR=0.001

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

fpn_c5p5	(Conv2D)
fpn_c4p4	(Conv2D)
fpn_c3p3	(Conv2D)
fpn_c2p2	(Conv2D)
fpn_p5	(Conv2D)
fpn_p2	(Conv2D)
fpn_p3	(Conv2D)
fpn_p4	(Conv2D)

In model: rpn_model

rpn_conv_shared	(Conv2D)
rpn_class_raw	(Conv2D)
rpn_bbox_pred	(Conv2D)
mrcnn_mask_conv1	(TimeDistributed)
mrcnn_mask_bn1	(TimeDistributed)
mrcnn_mask_conv2	(TimeDistributed)
mrcnn_mask_bn2	(TimeDistributed)
mrcnn_class_conv1	(TimeDistributed)
mrcnn_class_bn1	(TimeDistributed)
mrcnn_mask_conv3	(TimeDistributed)
mrcnn_mask_bn3	(TimeDistributed)
mrcnn_class_conv2	(TimeDistributed)
mrcnn_class_bn2	(TimeDistributed)
mrcnn_mask_conv4	(TimeDistributed)
mrcnn_mask_bn4	(TimeDistributed)
mrcnn_bbox_fc	(TimeDistributed)
mrcnn_mask_deconv	(TimeDistributed)
mrcnn_class_logits	(TimeDistributed)
mrcnn_mask	(TimeDistributed)

/home/paperspace/anaconda3/lib/python3.6/site-packages/tensorflow/python/ops/gradients_impl.py:100: UserWarning: Converting sparse IndexedSlices to a dense Tensor of unknown shape. This may consume a large amount of memory.

"Converting sparse IndexedSlices to a dense Tensor of unknown shape."

Epoch 1/1
 100/100 [=====] - 111s 1s/step - loss: 2.785
 9 - rpn_class_loss: 0.0184 - rpn_bbox_loss: 1.6422 - mrcnn_class_loss: 0.2197 - mrcnn_bbox_loss: 0.5045 - mrcnn_mask_loss: 0.4011 - val_loss: 1.9911 - val_rpn_class_loss: 0.0130 - val_rpn_bbox_loss: 1.2895 - val_mrcnn_class_loss: 0.0740 - val_mrcnn_bbox_loss: 0.3562 - val_mrcnn_mask_loss: 0.2584
 Epoch 0: Training batch 1 out of 40
 Re-starting from epoch 1

Starting at epoch 0. LR=0.00090000000000000001

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

fpn_c5p5	(Conv2D)
fpn_c4p4	(Conv2D)
fpn_c3p3	(Conv2D)
fpn_c2p2	(Conv2D)
fpn_p5	(Conv2D)
fpn_p2	(Conv2D)
fpn_p3	(Conv2D)
fpn_p4	(Conv2D)

In model: rpn_model

rpn_conv_shared	(Conv2D)
rpn_class_raw	(Conv2D)
rpn_bbox_pred	(Conv2D)
mrcnn_mask_conv1	(TimeDistributed)
mrcnn_mask_bn1	(TimeDistributed)
mrcnn_mask_conv2	(TimeDistributed)
mrcnn_mask_bn2	(TimeDistributed)
mrcnn_class_conv1	(TimeDistributed)
mrcnn_class_bn1	(TimeDistributed)
mrcnn_mask_conv3	(TimeDistributed)
mrcnn_mask_bn3	(TimeDistributed)
mrcnn_class_conv2	(TimeDistributed)
mrcnn_class_bn2	(TimeDistributed)
mrcnn_mask_conv4	(TimeDistributed)
mrcnn_mask_bn4	(TimeDistributed)
mrcnn_bbox_fc	(TimeDistributed)
mrcnn_mask_deconv	(TimeDistributed)
mrcnn_class_logits	(TimeDistributed)
mrcnn_mask	(TimeDistributed)

Epoch 1/1
 100/100 [=====] - 105s 1s/step - loss: 1.330
 5 - rpn_class_loss: 0.0151 - rpn_bbox_loss: 0.7286 - mrcnn_class_loss: 0.0598 - mrcnn_bbox_loss: 0.2592 - mrcnn_mask_loss: 0.2678 - val_loss: 2.2950 - val_rpn_class_loss: 0.0051 - val_rpn_bbox_loss: 1.5279 - val_mrcnn_class_loss: 0.0554 - val_mrcnn_bbox_loss: 0.3730 - val_mrcnn_mask_loss: 0.3335
 Epoch 0: Training batch 2 out of 40
 Re-starting from epoch 1

Starting at epoch 0. LR=0.00081000000000000001

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```
fpn_c5p5      (Conv2D)
fpn_c4p4      (Conv2D)
fpn_c3p3      (Conv2D)
fpn_c2p2      (Conv2D)
fpn_p5        (Conv2D)
fpn_p2        (Conv2D)
fpn_p3        (Conv2D)
fpn_p4        (Conv2D)
```

In model: rpn_model

```
    rpn_conv_shared      (Conv2D)
    rpn_class_raw        (Conv2D)
    rpn_bbox_pred        (Conv2D)
mrcnn_mask_conv1      (TimeDistributed)
mrcnn_mask_bn1        (TimeDistributed)
mrcnn_mask_conv2      (TimeDistributed)
mrcnn_mask_bn2        (TimeDistributed)
mrcnn_class_conv1      (TimeDistributed)
mrcnn_class_bn1        (TimeDistributed)
mrcnn_mask_conv3      (TimeDistributed)
mrcnn_mask_bn3        (TimeDistributed)
mrcnn_class_conv2      (TimeDistributed)
mrcnn_class_bn2        (TimeDistributed)
mrcnn_mask_conv4      (TimeDistributed)
mrcnn_mask_bn4        (TimeDistributed)
mrcnn_bbox_fc          (TimeDistributed)
mrcnn_mask_deconv      (TimeDistributed)
mrcnn_class_logits     (TimeDistributed)
mrcnn_mask             (TimeDistributed)
```

Epoch 1/1

100/100 [=====] - 101s 1s/step - loss: 1.083
 1 - rpn_class_loss: 0.0087 - rpn_bbox_loss: 0.6066 - mrcnn_class_loss: 0.0522 - mrcnn_bbox_loss: 0.1964 - mrcnn_mask_loss: 0.2191 - val_loss: 2.5264 - val_rpn_class_loss: 0.0174 - val_rpn_bbox_loss: 1.6280 - val_mrcnn_class_loss: 0.1139 - val_mrcnn_bbox_loss: 0.5140 - val_mrcnn_mask_loss: 0.2532

Epoch 0: Training batch 3 out of 40

Re-starting from epoch 1

Starting at epoch 0. LR=0.000729

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```
fpn_c5p5      (Conv2D)
fpn_c4p4      (Conv2D)
fpn_c3p3      (Conv2D)
fpn_c2p2      (Conv2D)
fpn_p5        (Conv2D)
fpn_p2        (Conv2D)
fpn_p3        (Conv2D)
fpn_p4        (Conv2D)
```

In model: rpn_model

```
    rpn_conv_shared      (Conv2D)
    rpn_class_raw        (Conv2D)
    rpn_bbox_pred        (Conv2D)
mrcnn_mask_conv1      (TimeDistributed)
```

```

mrcnn_mask_bn1          (TimeDistributed)
mrcnn_mask_conv2        (TimeDistributed)
mrcnn_mask_bn2          (TimeDistributed)
mrcnn_class_conv1       (TimeDistributed)
mrcnn_class_bn1         (TimeDistributed)
mrcnn_mask_conv3        (TimeDistributed)
mrcnn_mask_bn3          (TimeDistributed)
mrcnn_class_conv2       (TimeDistributed)
mrcnn_class_bn2         (TimeDistributed)
mrcnn_mask_conv4        (TimeDistributed)
mrcnn_mask_bn4          (TimeDistributed)
mrcnn_bbox_fc           (TimeDistributed)
mrcnn_mask_deconv       (TimeDistributed)
mrcnn_class_logits      (TimeDistributed)
mrcnn_mask              (TimeDistributed)

```

Epoch 1/1

100/100 [=====] - 101s 1s/step - loss: 1.211
 1 - rpn_class_loss: 0.0117 - rpn_bbox_loss: 0.7559 - mrcnn_class_loss: 0.0425 - mrcnn_bbox_loss: 0.1573 - mrcnn_mask_loss: 0.2438 - val_loss: 2.3194 - val_rpn_class_loss: 0.0105 - val_rpn_bbox_loss: 1.5851 - val_mrcnn_class_loss: 0.1528 - val_mrcnn_bbox_loss: 0.3942 - val_mrcnn_mask_loss: 0.1769

Epoch 0: Training batch 4 out of 40

Re-starting from epoch 1

Starting at epoch 0. LR=0.00065610000000000001

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```

fpn_c5p5          (Conv2D)
fpn_c4p4          (Conv2D)
fpn_c3p3          (Conv2D)
fpn_c2p2          (Conv2D)
fpn_p5            (Conv2D)
fpn_p2            (Conv2D)
fpn_p3            (Conv2D)
fpn_p4            (Conv2D)

```

In model: rpn_model

```

    rpn_conv_shared      (Conv2D)
    rpn_class_raw        (Conv2D)
    rpn_bbox_pred        (Conv2D)
mrcnn_mask_conv1       (TimeDistributed)
mrcnn_mask_bn1         (TimeDistributed)
mrcnn_mask_conv2       (TimeDistributed)
mrcnn_mask_bn2         (TimeDistributed)
mrcnn_class_conv1      (TimeDistributed)
mrcnn_class_bn1        (TimeDistributed)
mrcnn_mask_conv3       (TimeDistributed)
mrcnn_mask_bn3         (TimeDistributed)
mrcnn_class_conv2      (TimeDistributed)
mrcnn_class_bn2        (TimeDistributed)
mrcnn_mask_conv4       (TimeDistributed)
mrcnn_mask_bn4         (TimeDistributed)
mrcnn_bbox_fc          (TimeDistributed)
mrcnn_mask_deconv      (TimeDistributed)
mrcnn_class_logits     (TimeDistributed)

```

```

mrcnn_mask (TimeDistributed)
Epoch 1/1
100/100 [=====] - 105s 1s/step - loss: 0.760
1 - rpn_class_loss: 0.0066 - rpn_bbox_loss: 0.3075 - mrcnn_class_loss: 0.0477 - mrcnn_bbox_loss: 0.1107 - mrcnn_mask_loss: 0.2876 - val_loss: 1.3385 - val_rpn_class_loss: 0.0037 - val_rpn_bbox_loss: 0.8166 - val_mrcnn_class_loss: 0.0713 - val_mrcnn_bbox_loss: 0.2829 - val_mrcnn_mask_loss: 0.1639
Epoch 0: Training batch 5 out of 40
Re-starting from epoch 1

```

Starting at epoch 0. LR=0.00059049

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```

fpn_c5p5 (Conv2D)
fpn_c4p4 (Conv2D)
fpn_c3p3 (Conv2D)
fpn_c2p2 (Conv2D)
fpn_p5 (Conv2D)
fpn_p2 (Conv2D)
fpn_p3 (Conv2D)
fpn_p4 (Conv2D)

```

In model: rpn_model

```

    rpn_conv_shared (Conv2D)
    rpn_class_raw (Conv2D)
    rpn_bbox_pred (Conv2D)
mrcnn_mask_conv1 (TimeDistributed)
mrcnn_mask_bn1 (TimeDistributed)
mrcnn_mask_conv2 (TimeDistributed)
mrcnn_mask_bn2 (TimeDistributed)
mrcnn_class_conv1 (TimeDistributed)
mrcnn_class_bn1 (TimeDistributed)
mrcnn_mask_conv3 (TimeDistributed)
mrcnn_mask_bn3 (TimeDistributed)
mrcnn_class_conv2 (TimeDistributed)
mrcnn_class_bn2 (TimeDistributed)
mrcnn_mask_conv4 (TimeDistributed)
mrcnn_mask_bn4 (TimeDistributed)
mrcnn_bbox_fc (TimeDistributed)
mrcnn_mask_deconv (TimeDistributed)
mrcnn_class_logits (TimeDistributed)
mrcnn_mask (TimeDistributed)

```

```

Epoch 1/1
100/100 [=====] - 102s 1s/step - loss: 1.106
8 - rpn_class_loss: 0.0083 - rpn_bbox_loss: 0.6501 - mrcnn_class_loss: 0.0653 - mrcnn_bbox_loss: 0.1591 - mrcnn_mask_loss: 0.2240 - val_loss: 2.2187 - val_rpn_class_loss: 0.0064 - val_rpn_bbox_loss: 1.4020 - val_mrcnn_class_loss: 0.1336 - val_mrcnn_bbox_loss: 0.4555 - val_mrcnn_mask_loss: 0.2212
Epoch 0: Training batch 6 out of 40
Re-starting from epoch 1

```

Starting at epoch 0. LR=0.000531441

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1

835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```
fpn_c5p5          (Conv2D)
fpn_c4p4          (Conv2D)
fpn_c3p3          (Conv2D)
fpn_c2p2          (Conv2D)
fpn_p5            (Conv2D)
fpn_p2            (Conv2D)
fpn_p3            (Conv2D)
fpn_p4            (Conv2D)
```

In model: rpn_model

```
    rpn_conv_shared      (Conv2D)
    rpn_class_raw         (Conv2D)
    rpn_bbox_pred         (Conv2D)
mrcnn_mask_conv1        (TimeDistributed)
mrcnn_mask_bn1          (TimeDistributed)
mrcnn_mask_conv2        (TimeDistributed)
mrcnn_mask_bn2          (TimeDistributed)
mrcnn_class_conv1       (TimeDistributed)
mrcnn_class_bn1         (TimeDistributed)
mrcnn_mask_conv3        (TimeDistributed)
mrcnn_mask_bn3          (TimeDistributed)
mrcnn_class_conv2       (TimeDistributed)
mrcnn_class_bn2         (TimeDistributed)
mrcnn_mask_conv4        (TimeDistributed)
mrcnn_mask_bn4          (TimeDistributed)
mrcnn_bbox_fc           (TimeDistributed)
mrcnn_mask_deconv       (TimeDistributed)
mrcnn_class_logits      (TimeDistributed)
mrcnn_mask              (TimeDistributed)
```

Epoch 1/1

100/100 [=====] - 102s 1s/step - loss: 1.0559 - rpn_class_loss: 0.0075 - rpn_bbox_loss: 0.5578 - mrcnn_class_loss: 0.0813 - mrcnn_bbox_loss: 0.1656 - mrcnn_mask_loss: 0.2437 - val_loss: 2.0538 - val_rpn_class_loss: 0.0042 - val_rpn_bbox_loss: 1.3202 - val_mrcnn_class_loss: 0.0902 - val_mrcnn_bbox_loss: 0.3837 - val_mrcnn_mask_loss: 0.2556

Epoch 0: Training batch 7 out of 40

Re-starting from epoch 1

Starting at epoch 0. LR=0.0004782969

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1

835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```
fpn_c5p5          (Conv2D)
fpn_c4p4          (Conv2D)
fpn_c3p3          (Conv2D)
fpn_c2p2          (Conv2D)
fpn_p5            (Conv2D)
fpn_p2            (Conv2D)
fpn_p3            (Conv2D)
fpn_p4            (Conv2D)
```

In model: rpn_model

```
    rpn_conv_shared      (Conv2D)
    rpn_class_raw         (Conv2D)
    rpn_bbox_pred         (Conv2D)
```

```

mrcnn_mask_conv1      (TimeDistributed)
mrcnn_mask_bn1        (TimeDistributed)
mrcnn_mask_conv2      (TimeDistributed)
mrcnn_mask_bn2        (TimeDistributed)
mrcnn_class_conv1     (TimeDistributed)
mrcnn_class_bn1       (TimeDistributed)
mrcnn_mask_conv3      (TimeDistributed)
mrcnn_mask_bn3        (TimeDistributed)
mrcnn_class_conv2     (TimeDistributed)
mrcnn_class_bn2       (TimeDistributed)
mrcnn_mask_conv4      (TimeDistributed)
mrcnn_mask_bn4        (TimeDistributed)
mrcnn_bbox_fc         (TimeDistributed)
mrcnn_mask_deconv     (TimeDistributed)
mrcnn_class_logits    (TimeDistributed)
mrcnn_mask            (TimeDistributed)

```

Epoch 1/1

100/100 [=====] - 102s 1s/step - loss: 1.289
 0 - rpn_class_loss: 0.0052 - rpn_bbox_loss: 0.7337 - mrcnn_class_loss: 0.0909 - mrcnn_bbox_loss: 0.1743 - mrcnn_mask_loss: 0.2850 - val_loss: 2.2413 - val_rpn_class_loss: 0.0051 - val_rpn_bbox_loss: 1.5917 - val_mrcnn_class_loss: 0.0846 - val_mrcnn_bbox_loss: 0.3105 - val_mrcnn_mask_loss: 0.2493

Epoch 0: Training batch 8 out of 40

Re-starting from epoch 1

Starting at epoch 0. LR=0.00043046721

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```

fpn_c5p5              (Conv2D)
fpn_c4p4              (Conv2D)
fpn_c3p3              (Conv2D)
fpn_c2p2              (Conv2D)
fpn_p5                (Conv2D)
fpn_p2                (Conv2D)
fpn_p3                (Conv2D)
fpn_p4                (Conv2D)

```

In model: rpn_model

```

    rpn_conv_shared    (Conv2D)
    rpn_class_raw       (Conv2D)
    rpn_bbox_pred       (Conv2D)
mrcnn_mask_conv1      (TimeDistributed)
mrcnn_mask_bn1        (TimeDistributed)
mrcnn_mask_conv2      (TimeDistributed)
mrcnn_mask_bn2        (TimeDistributed)
mrcnn_class_conv1     (TimeDistributed)
mrcnn_class_bn1       (TimeDistributed)
mrcnn_mask_conv3      (TimeDistributed)
mrcnn_mask_bn3        (TimeDistributed)
mrcnn_class_conv2     (TimeDistributed)
mrcnn_class_bn2       (TimeDistributed)
mrcnn_mask_conv4      (TimeDistributed)
mrcnn_mask_bn4        (TimeDistributed)
mrcnn_bbox_fc         (TimeDistributed)
mrcnn_mask_deconv     (TimeDistributed)

```



```

mrcnn_class_logits      (TimeDistributed)
mrcnn_mask               (TimeDistributed)
Epoch 1/1
100/100 [=====] - 102s 1s/step - loss: 1.182
9 - rpn_class_loss: 0.0069 - rpn_bbox_loss: 0.7495 - mrcnn_class_loss: 0.0728 - mrcnn_bbox_loss: 0.1570 - mrcnn_mask_loss: 0.1967 - val_loss: 1.8006 - val_rpn_class_loss: 0.0048 - val_rpn_bbox_loss: 1.0078 - val_mrcnn_class_loss: 0.1290 - val_mrcnn_bbox_loss: 0.4357 - val_mrcnn_mask_loss: 0.2233
Epoch 0: Training batch 9 out of 40
Re-starting from epoch 1

```

Starting at epoch 0. LR=0.000387420489

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```

fpn_c5p5      (Conv2D)
fpn_c4p4      (Conv2D)
fpn_c3p3      (Conv2D)
fpn_c2p2      (Conv2D)
fpn_p5        (Conv2D)
fpn_p2        (Conv2D)
fpn_p3        (Conv2D)
fpn_p4        (Conv2D)

```

In model: rpn_model

```

    rpn_conv_shared      (Conv2D)
    rpn_class_raw         (Conv2D)
    rpn_bbox_pred         (Conv2D)
mrcnn_mask_conv1        (TimeDistributed)
mrcnn_mask_bn1          (TimeDistributed)
mrcnn_mask_conv2        (TimeDistributed)
mrcnn_mask_bn2          (TimeDistributed)
mrcnn_class_conv1       (TimeDistributed)
mrcnn_class_bn1         (TimeDistributed)
mrcnn_mask_conv3        (TimeDistributed)
mrcnn_mask_bn3          (TimeDistributed)
mrcnn_class_conv2       (TimeDistributed)
mrcnn_class_bn2         (TimeDistributed)
mrcnn_mask_conv4        (TimeDistributed)
mrcnn_mask_bn4          (TimeDistributed)
mrcnn_bbox_fc           (TimeDistributed)
mrcnn_mask_deconv       (TimeDistributed)
mrcnn_class_logits      (TimeDistributed)
mrcnn_mask               (TimeDistributed)

```

```

Epoch 1/1
100/100 [=====] - 103s 1s/step - loss: 0.968
1 - rpn_class_loss: 0.0048 - rpn_bbox_loss: 0.5819 - mrcnn_class_loss: 0.0636 - mrcnn_bbox_loss: 0.1300 - mrcnn_mask_loss: 0.1878 - val_loss: 2.0970 - val_rpn_class_loss: 0.0079 - val_rpn_bbox_loss: 1.4279 - val_mrcnn_class_loss: 0.0921 - val_mrcnn_bbox_loss: 0.3069 - val_mrcnn_mask_loss: 0.2621
Epoch 0: Training batch 10 out of 40
Re-starting from epoch 1

```

Starting at epoch 0. LR=0.0003486784401

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1
835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```
fpn_c5p5          (Conv2D)
fpn_c4p4          (Conv2D)
fpn_c3p3          (Conv2D)
fpn_c2p2          (Conv2D)
fpn_p5            (Conv2D)
fpn_p2            (Conv2D)
fpn_p3            (Conv2D)
fpn_p4            (Conv2D)
```

In model: rpn_model

```
    rpn_conv_shared      (Conv2D)
    rpn_class_raw        (Conv2D)
    rpn_bbox_pred        (Conv2D)
mrcnn_mask_conv1      (TimeDistributed)
mrcnn_mask_bn1        (TimeDistributed)
mrcnn_mask_conv2      (TimeDistributed)
mrcnn_mask_bn2        (TimeDistributed)
mrcnn_class_conv1     (TimeDistributed)
mrcnn_class_bn1       (TimeDistributed)
mrcnn_mask_conv3      (TimeDistributed)
mrcnn_mask_bn3        (TimeDistributed)
mrcnn_class_conv2     (TimeDistributed)
mrcnn_class_bn2       (TimeDistributed)
mrcnn_mask_conv4      (TimeDistributed)
mrcnn_mask_bn4        (TimeDistributed)
mrcnn_bbox_fc         (TimeDistributed)
mrcnn_mask_deconv     (TimeDistributed)
mrcnn_class_logits    (TimeDistributed)
mrcnn_mask            (TimeDistributed)
```

Epoch 1/1

100/100 [=====] - 101s 1s/step - loss: 0.924
0 - rpn_class_loss: 0.0036 - rpn_bbox_loss: 0.5368 - mrcnn_class_loss: 0.0471 - mrcnn_bbox_loss: 0.1306 - mrcnn_mask_loss: 0.2059 - val_loss: 1.1579 - val_rpn_class_loss: 0.0046 - val_rpn_bbox_loss: 0.7345 - val_mrcnn_class_loss: 0.0465 - val_mrcnn_bbox_loss: 0.1494 - val_mrcnn_mask_loss: 0.2228

Epoch 0: Training batch 11 out of 40

Re-starting from epoch 1

Starting at epoch 0. LR=0.00031381059609000004

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1
835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```
fpn_c5p5          (Conv2D)
fpn_c4p4          (Conv2D)
fpn_c3p3          (Conv2D)
fpn_c2p2          (Conv2D)
fpn_p5            (Conv2D)
fpn_p2            (Conv2D)
fpn_p3            (Conv2D)
fpn_p4            (Conv2D)
```

In model: rpn_model

```
    rpn_conv_shared      (Conv2D)
    rpn_class_raw        (Conv2D)
```

rpn_bbox_pred	(Conv2D)
mrcnn_mask_conv1	(TimeDistributed)
mrcnn_mask_bn1	(TimeDistributed)
mrcnn_mask_conv2	(TimeDistributed)
mrcnn_mask_bn2	(TimeDistributed)
mrcnn_class_conv1	(TimeDistributed)
mrcnn_class_bn1	(TimeDistributed)
mrcnn_mask_conv3	(TimeDistributed)
mrcnn_mask_bn3	(TimeDistributed)
mrcnn_class_conv2	(TimeDistributed)
mrcnn_class_bn2	(TimeDistributed)
mrcnn_mask_conv4	(TimeDistributed)
mrcnn_mask_bn4	(TimeDistributed)
mrcnn_bbox_fc	(TimeDistributed)
mrcnn_mask_deconv	(TimeDistributed)
mrcnn_class_logits	(TimeDistributed)
mrcnn_mask	(TimeDistributed)

Epoch 1/1

100/100 [=====] - 103s 1s/step - loss: 1.127
 8 - rpn_class_loss: 0.0053 - rpn_bbox_loss: 0.7540 - mrcnn_class_loss: 0.0651 - mrcnn_bbox_loss: 0.1647 - mrcnn_mask_loss: 0.1387 - val_loss: 2.2728 - val_rpn_class_loss: 0.0063 - val_rpn_bbox_loss: 1.4662 - val_mrcnn_class_loss: 0.1221 - val_mrcnn_bbox_loss: 0.4408 - val_mrcnn_mask_loss: 0.2373

Epoch 0: Training batch 12 out of 40

Re-starting from epoch 1

Starting at epoch 0. LR=0.00028242953648100003

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

fpn_c5p5	(Conv2D)
fpn_c4p4	(Conv2D)
fpn_c3p3	(Conv2D)
fpn_c2p2	(Conv2D)
fpn_p5	(Conv2D)
fpn_p2	(Conv2D)
fpn_p3	(Conv2D)
fpn_p4	(Conv2D)

In model: rpn_model

rpn_conv_shared	(Conv2D)
rpn_class_raw	(Conv2D)
rpn_bbox_pred	(Conv2D)
mrcnn_mask_conv1	(TimeDistributed)
mrcnn_mask_bn1	(TimeDistributed)
mrcnn_mask_conv2	(TimeDistributed)
mrcnn_mask_bn2	(TimeDistributed)
mrcnn_class_conv1	(TimeDistributed)
mrcnn_class_bn1	(TimeDistributed)
mrcnn_mask_conv3	(TimeDistributed)
mrcnn_mask_bn3	(TimeDistributed)
mrcnn_class_conv2	(TimeDistributed)
mrcnn_class_bn2	(TimeDistributed)
mrcnn_mask_conv4	(TimeDistributed)
mrcnn_mask_bn4	(TimeDistributed)
mrcnn_bbox_fc	(TimeDistributed)

```

mrcnn_mask_deconv      (TimeDistributed)
mrcnn_class_logits      (TimeDistributed)
mrcnn_mask              (TimeDistributed)
Epoch 1/1
100/100 [=====] - 104s 1s/step - loss: 0.671
7 - rpn_class_loss: 0.0038 - rpn_bbox_loss: 0.3141 - mrcnn_class_loss: 0.0454 - mrcnn_bbox_loss: 0.0984 - mrcnn_mask_loss: 0.2100 - val_loss: 1.6837 - val_rpn_class_loss: 0.0053 - val_rpn_bbox_loss: 1.0427 - val_mrcnn_class_loss: 0.1356 - val_mrcnn_bbox_loss: 0.3423 - val_mrcnn_mask_loss: 0.1578
Epoch 0: Training batch 13 out of 40
Re-starting from epoch 1

```

Starting at epoch 0. LR=0.00025418658283290005

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```

fpn_c5p5      (Conv2D)
fpn_c4p4      (Conv2D)
fpn_c3p3      (Conv2D)
fpn_c2p2      (Conv2D)
fpn_p5        (Conv2D)
fpn_p2        (Conv2D)
fpn_p3        (Conv2D)
fpn_p4        (Conv2D)

```

In model: rpn_model

```

    rpn_conv_shared      (Conv2D)
    rpn_class_raw        (Conv2D)
    rpn_bbox_pred        (Conv2D)
mrcnn_mask_conv1      (TimeDistributed)
mrcnn_mask_bn1        (TimeDistributed)
mrcnn_mask_conv2      (TimeDistributed)
mrcnn_mask_bn2        (TimeDistributed)
mrcnn_class_conv1     (TimeDistributed)
mrcnn_class_bn1       (TimeDistributed)
mrcnn_mask_conv3      (TimeDistributed)
mrcnn_mask_bn3        (TimeDistributed)
mrcnn_class_conv2     (TimeDistributed)
mrcnn_class_bn2       (TimeDistributed)
mrcnn_mask_conv4      (TimeDistributed)
mrcnn_mask_bn4        (TimeDistributed)
mrcnn_bbox_fc         (TimeDistributed)
mrcnn_mask_deconv     (TimeDistributed)
mrcnn_class_logits     (TimeDistributed)
mrcnn_mask            (TimeDistributed)

```

```

Epoch 1/1
100/100 [=====] - 102s 1s/step - loss: 0.888
2 - rpn_class_loss: 0.0047 - rpn_bbox_loss: 0.4392 - mrcnn_class_loss: 0.0585 - mrcnn_bbox_loss: 0.1537 - mrcnn_mask_loss: 0.2322 - val_loss: 1.4106 - val_rpn_class_loss: 0.0023 - val_rpn_bbox_loss: 0.9446 - val_mrcnn_class_loss: 0.0393 - val_mrcnn_bbox_loss: 0.2372 - val_mrcnn_mask_loss: 0.1873
Epoch 0: Training batch 14 out of 40
Re-starting from epoch 1

```

Starting at epoch 0. LR=0.00022876792454961005

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```
fpn_c5p5      (Conv2D)
fpn_c4p4      (Conv2D)
fpn_c3p3      (Conv2D)
fpn_c2p2      (Conv2D)
fpn_p5        (Conv2D)
fpn_p2        (Conv2D)
fpn_p3        (Conv2D)
fpn_p4        (Conv2D)
```

In model: rpn_model

```
    rpn_conv_shared      (Conv2D)
    rpn_class_raw         (Conv2D)
    rpn_bbox_pred         (Conv2D)
mrcnn_mask_conv1        (TimeDistributed)
mrcnn_mask_bn1          (TimeDistributed)
mrcnn_mask_conv2        (TimeDistributed)
mrcnn_mask_bn2          (TimeDistributed)
mrcnn_class_conv1       (TimeDistributed)
mrcnn_class_bn1         (TimeDistributed)
mrcnn_mask_conv3        (TimeDistributed)
mrcnn_mask_bn3          (TimeDistributed)
mrcnn_class_conv2       (TimeDistributed)
mrcnn_class_bn2         (TimeDistributed)
mrcnn_mask_conv4        (TimeDistributed)
mrcnn_mask_bn4          (TimeDistributed)
mrcnn_bbox_fc           (TimeDistributed)
mrcnn_mask_deconv       (TimeDistributed)
mrcnn_class_logits      (TimeDistributed)
mrcnn_mask              (TimeDistributed)
```

Epoch 1/1

100/100 [=====] - 103s 1s/step - loss: 0.798
 2 - rpn_class_loss: 0.0040 - rpn_bbox_loss: 0.4177 - mrcnn_class_loss: 0.0479 - mrcnn_bbox_loss: 0.1370 - mrcnn_mask_loss: 0.1915 - val_loss: 1.6841 - val_rpn_class_loss: 0.0057 - val_rpn_bbox_loss: 1.1825 - val_mrcnn_class_loss: 0.1030 - val_mrcnn_bbox_loss: 0.2261 - val_mrcnn_mask_loss: 0.1668

Epoch 0: Training batch 15 out of 40

Re-starting from epoch 1

Starting at epoch 0. LR=0.00020589113209464906

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```
fpn_c5p5      (Conv2D)
fpn_c4p4      (Conv2D)
fpn_c3p3      (Conv2D)
fpn_c2p2      (Conv2D)
fpn_p5        (Conv2D)
fpn_p2        (Conv2D)
fpn_p3        (Conv2D)
fpn_p4        (Conv2D)
```

In model: rpn_model

```
    rpn_conv_shared      (Conv2D)
```

rpn_class_raw	(Conv2D)
rpn_bbox_pred	(Conv2D)
mrcnn_mask_conv1	(TimeDistributed)
mrcnn_mask_bn1	(TimeDistributed)
mrcnn_mask_conv2	(TimeDistributed)
mrcnn_mask_bn2	(TimeDistributed)
mrcnn_class_conv1	(TimeDistributed)
mrcnn_class_bn1	(TimeDistributed)
mrcnn_mask_conv3	(TimeDistributed)
mrcnn_mask_bn3	(TimeDistributed)
mrcnn_class_conv2	(TimeDistributed)
mrcnn_class_bn2	(TimeDistributed)
mrcnn_mask_conv4	(TimeDistributed)
mrcnn_mask_bn4	(TimeDistributed)
mrcnn_bbox_fc	(TimeDistributed)
mrcnn_mask_deconv	(TimeDistributed)
mrcnn_class_logits	(TimeDistributed)
mrcnn_mask	(TimeDistributed)

Epoch 1/1

100/100 [=====] - 104s 1s/step - loss: 0.8507 - rpn_class_loss: 0.0049 - rpn_bbox_loss: 0.4889 - mrcnn_class_loss: 0.0535 - mrcnn_bbox_loss: 0.1215 - mrcnn_mask_loss: 0.1818 - val_loss: 1.3925 - val_rpn_class_loss: 0.0031 - val_rpn_bbox_loss: 0.9574 - val_mrcnn_class_loss: 0.0445 - val_mrcnn_bbox_loss: 0.1888 - val_mrcnn_mask_loss: 0.1987

Epoch 0: Training batch 16 out of 40

Re-starting from epoch 1

Starting at epoch 0. LR=0.00018530201888518417

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

fpn_c5p5	(Conv2D)
fpn_c4p4	(Conv2D)
fpn_c3p3	(Conv2D)
fpn_c2p2	(Conv2D)
fpn_p5	(Conv2D)
fpn_p2	(Conv2D)
fpn_p3	(Conv2D)
fpn_p4	(Conv2D)

In model: rpn_model

rpn_conv_shared	(Conv2D)
rpn_class_raw	(Conv2D)
rpn_bbox_pred	(Conv2D)
mrcnn_mask_conv1	(TimeDistributed)
mrcnn_mask_bn1	(TimeDistributed)
mrcnn_mask_conv2	(TimeDistributed)
mrcnn_mask_bn2	(TimeDistributed)
mrcnn_class_conv1	(TimeDistributed)
mrcnn_class_bn1	(TimeDistributed)
mrcnn_mask_conv3	(TimeDistributed)
mrcnn_mask_bn3	(TimeDistributed)
mrcnn_class_conv2	(TimeDistributed)
mrcnn_class_bn2	(TimeDistributed)
mrcnn_mask_conv4	(TimeDistributed)
mrcnn_mask_bn4	(TimeDistributed)

```

mrcnn_bbox_fc          (TimeDistributed)
mrcnn_mask_deconv      (TimeDistributed)
mrcnn_class_logits     (TimeDistributed)
mrcnn_mask              (TimeDistributed)
Epoch 1/1
100/100 [=====] - 104s 1s/step - loss: 1.0302 - rpn_class_loss: 0.0044 - rpn_bbox_loss: 0.6393 - mrcnn_class_loss: 0.0537 - mrcnn_bbox_loss: 0.1373 - mrcnn_mask_loss: 0.1955 - val_loss: 1.6983 - val_rpn_class_loss: 0.0034 - val_rpn_bbox_loss: 1.0574 - val_mrcnn_class_loss: 0.0905 - val_mrcnn_bbox_loss: 0.2787 - val_mrcnn_mask_loss: 0.2684
Epoch 0: Training batch 17 out of 40
Re-starting from epoch 1

```

Starting at epoch 0. LR=0.00016677181699666576

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```

fpn_c5p5          (Conv2D)
fpn_c4p4          (Conv2D)
fpn_c3p3          (Conv2D)
fpn_c2p2          (Conv2D)
fpn_p5            (Conv2D)
fpn_p2            (Conv2D)
fpn_p3            (Conv2D)
fpn_p4            (Conv2D)

```

In model: rpn_model

```

    rpn_conv_shared      (Conv2D)
    rpn_class_raw         (Conv2D)
    rpn_bbox_pred         (Conv2D)
mrcnn_mask_conv1        (TimeDistributed)
mrcnn_mask_bn1          (TimeDistributed)
mrcnn_mask_conv2        (TimeDistributed)
mrcnn_mask_bn2          (TimeDistributed)
mrcnn_class_conv1       (TimeDistributed)
mrcnn_class_bn1         (TimeDistributed)
mrcnn_mask_conv3        (TimeDistributed)
mrcnn_mask_bn3          (TimeDistributed)
mrcnn_class_conv2       (TimeDistributed)
mrcnn_class_bn2         (TimeDistributed)
mrcnn_mask_conv4        (TimeDistributed)
mrcnn_mask_bn4          (TimeDistributed)
mrcnn_bbox_fc           (TimeDistributed)
mrcnn_mask_deconv       (TimeDistributed)
mrcnn_class_logits      (TimeDistributed)
mrcnn_mask              (TimeDistributed)

```

```

Epoch 1/1
100/100 [=====] - 104s 1s/step - loss: 0.9509 - rpn_class_loss: 0.0039 - rpn_bbox_loss: 0.5277 - mrcnn_class_loss: 0.0452 - mrcnn_bbox_loss: 0.1260 - mrcnn_mask_loss: 0.2481 - val_loss: 1.3124 - val_rpn_class_loss: 0.0034 - val_rpn_bbox_loss: 0.7526 - val_mrcnn_class_loss: 0.0516 - val_mrcnn_bbox_loss: 0.2609 - val_mrcnn_mask_loss: 0.2438
Epoch 0: Training batch 18 out of 40
Re-starting from epoch 1

```

Starting at epoch 0. LR=0.0001500946352969992

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```
fpn_c5p5          (Conv2D)
fpn_c4p4          (Conv2D)
fpn_c3p3          (Conv2D)
fpn_c2p2          (Conv2D)
fpn_p5            (Conv2D)
fpn_p2            (Conv2D)
fpn_p3            (Conv2D)
fpn_p4            (Conv2D)
```

In model: rpn_model

```
    rpn_conv_shared      (Conv2D)
    rpn_class_raw        (Conv2D)
    rpn_bbox_pred        (Conv2D)
mrcnn_mask_conv1      (TimeDistributed)
mrcnn_mask_bn1        (TimeDistributed)
mrcnn_mask_conv2      (TimeDistributed)
mrcnn_mask_bn2        (TimeDistributed)
mrcnn_class_conv1     (TimeDistributed)
mrcnn_class_bn1       (TimeDistributed)
mrcnn_mask_conv3      (TimeDistributed)
mrcnn_mask_bn3        (TimeDistributed)
mrcnn_class_conv2     (TimeDistributed)
mrcnn_class_bn2       (TimeDistributed)
mrcnn_mask_conv4      (TimeDistributed)
mrcnn_mask_bn4        (TimeDistributed)
mrcnn_bbox_fc         (TimeDistributed)
mrcnn_mask_deconv     (TimeDistributed)
mrcnn_class_logits    (TimeDistributed)
mrcnn_mask            (TimeDistributed)
```

Epoch 1/1

100/100 [=====] - 103s 1s/step - loss: 1.170
 0 - rpn_class_loss: 0.0054 - rpn_bbox_loss: 0.7688 - mrcnn_class_loss: 0.0536 - mrcnn_bbox_loss: 0.1187 - mrcnn_mask_loss: 0.2235 - val_loss: 2.1357 - val_rpn_class_loss: 0.0061 - val_rpn_bbox_loss: 1.4862 - val_mrcnn_class_loss: 0.1555 - val_mrcnn_bbox_loss: 0.2789 - val_mrcnn_mask_loss: 0.2090

Epoch 0: Training batch 19 out of 40

Re-starting from epoch 1

Starting at epoch 0. LR=0.0001350851717672993

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```
fpn_c5p5          (Conv2D)
fpn_c4p4          (Conv2D)
fpn_c3p3          (Conv2D)
fpn_c2p2          (Conv2D)
fpn_p5            (Conv2D)
fpn_p2            (Conv2D)
fpn_p3            (Conv2D)
fpn_p4            (Conv2D)
```

In model: rpn_model

rpn_conv_shared	(Conv2D)
rpn_class_raw	(Conv2D)
rpn_bbox_pred	(Conv2D)
mrcnn_mask_conv1	(TimeDistributed)
mrcnn_mask_bn1	(TimeDistributed)
mrcnn_mask_conv2	(TimeDistributed)
mrcnn_mask_bn2	(TimeDistributed)
mrcnn_class_conv1	(TimeDistributed)
mrcnn_class_bn1	(TimeDistributed)
mrcnn_mask_conv3	(TimeDistributed)
mrcnn_mask_bn3	(TimeDistributed)
mrcnn_class_conv2	(TimeDistributed)
mrcnn_class_bn2	(TimeDistributed)
mrcnn_mask_conv4	(TimeDistributed)
mrcnn_mask_bn4	(TimeDistributed)
mrcnn_bbox_fc	(TimeDistributed)
mrcnn_mask_deconv	(TimeDistributed)
mrcnn_class_logits	(TimeDistributed)
mrcnn_mask	(TimeDistributed)

Epoch 1/1

100/100 [=====] - 104s 1s/step - loss: 0.7683 - rpn_class_loss: 0.0045 - rpn_bbox_loss: 0.3980 - mrcnn_class_loss: 0.0444 - mrcnn_bbox_loss: 0.1144 - mrcnn_mask_loss: 0.2069 - val_loss: 2.2610 - val_rpn_class_loss: 0.0044 - val_rpn_bbox_loss: 1.5530 - val_mrcnn_class_loss: 0.0882 - val_mrcnn_bbox_loss: 0.3971 - val_mrcnn_mask_loss: 0.2182

Epoch 0: Training batch 20 out of 40

Re-starting from epoch 1

Starting at epoch 0. LR=0.00012157665459056936

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

fpn_c5p5	(Conv2D)
fpn_c4p4	(Conv2D)
fpn_c3p3	(Conv2D)
fpn_c2p2	(Conv2D)
fpn_p5	(Conv2D)
fpn_p2	(Conv2D)
fpn_p3	(Conv2D)
fpn_p4	(Conv2D)

In model: rpn_model

rpn_conv_shared	(Conv2D)
rpn_class_raw	(Conv2D)
rpn_bbox_pred	(Conv2D)
mrcnn_mask_conv1	(TimeDistributed)
mrcnn_mask_bn1	(TimeDistributed)
mrcnn_mask_conv2	(TimeDistributed)
mrcnn_mask_bn2	(TimeDistributed)
mrcnn_class_conv1	(TimeDistributed)
mrcnn_class_bn1	(TimeDistributed)
mrcnn_mask_conv3	(TimeDistributed)
mrcnn_mask_bn3	(TimeDistributed)
mrcnn_class_conv2	(TimeDistributed)
mrcnn_class_bn2	(TimeDistributed)
mrcnn_mask_conv4	(TimeDistributed)

```

mrcnn_mask_bn4          (TimeDistributed)
mrcnn_bbox_fc           (TimeDistributed)
mrcnn_mask_deconv       (TimeDistributed)
mrcnn_class_logits      (TimeDistributed)
mrcnn_mask              (TimeDistributed)
Epoch 1/1
100/100 [=====] - 103s 1s/step - loss: 1.053
8 - rpn_class_loss: 0.0046 - rpn_bbox_loss: 0.5862 - mrcnn_class_loss: 0.0476 - mrcnn_bbox_loss: 0.1650 - mrcnn_mask_loss: 0.2505 - val_loss: 1.1941 - val_rpn_class_loss: 0.0026 - val_rpn_bbox_loss: 0.7229 - val_mrcnn_class_loss: 0.0360 - val_mrcnn_bbox_loss: 0.2442 - val_mrcnn_mask_loss: 0.1884
Epoch 0: Training batch 21 out of 40
Re-starting from epoch 1

```

Starting at epoch 0. LR=0.00010941898913151243

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```

fpn_c5p5          (Conv2D)
fpn_c4p4          (Conv2D)
fpn_c3p3          (Conv2D)
fpn_c2p2          (Conv2D)
fpn_p5            (Conv2D)
fpn_p2            (Conv2D)
fpn_p3            (Conv2D)
fpn_p4            (Conv2D)

```

In model: rpn_model

```

    rpn_conv_shared      (Conv2D)
    rpn_class_raw        (Conv2D)
    rpn_bbox_pred        (Conv2D)
mrcnn_mask_conv1       (TimeDistributed)
mrcnn_mask_bn1         (TimeDistributed)
mrcnn_mask_conv2       (TimeDistributed)
mrcnn_mask_bn2         (TimeDistributed)
mrcnn_class_conv1      (TimeDistributed)
mrcnn_class_bn1        (TimeDistributed)
mrcnn_mask_conv3       (TimeDistributed)
mrcnn_mask_bn3         (TimeDistributed)
mrcnn_class_conv2      (TimeDistributed)
mrcnn_class_bn2        (TimeDistributed)
mrcnn_mask_conv4       (TimeDistributed)
mrcnn_mask_bn4         (TimeDistributed)
mrcnn_bbox_fc          (TimeDistributed)
mrcnn_mask_deconv      (TimeDistributed)
mrcnn_class_logits     (TimeDistributed)
mrcnn_mask             (TimeDistributed)

```

```

Epoch 1/1
100/100 [=====] - 105s 1s/step - loss: 0.885
3 - rpn_class_loss: 0.0032 - rpn_bbox_loss: 0.5354 - mrcnn_class_loss: 0.0374 - mrcnn_bbox_loss: 0.1121 - mrcnn_mask_loss: 0.1973 - val_loss: 1.3984 - val_rpn_class_loss: 0.0052 - val_rpn_bbox_loss: 0.8502 - val_mrcnn_class_loss: 0.0857 - val_mrcnn_bbox_loss: 0.2623 - val_mrcnn_mask_loss: 0.1949
Epoch 0: Training batch 22 out of 40
Re-starting from epoch 1

```

Starting at epoch 0. LR=9.847709021836118e-05

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```
fpn_c5p5          (Conv2D)
fpn_c4p4          (Conv2D)
fpn_c3p3          (Conv2D)
fpn_c2p2          (Conv2D)
fpn_p5            (Conv2D)
fpn_p2            (Conv2D)
fpn_p3            (Conv2D)
fpn_p4            (Conv2D)
```

In model: rpn_model

```
    rpn_conv_shared (Conv2D)
    rpn_class_raw    (Conv2D)
    rpn_bbox_pred    (Conv2D)
mrcnn_mask_conv1    (TimeDistributed)
mrcnn_mask_bn1      (TimeDistributed)
mrcnn_mask_conv2    (TimeDistributed)
mrcnn_mask_bn2      (TimeDistributed)
mrcnn_class_conv1   (TimeDistributed)
mrcnn_class_bn1     (TimeDistributed)
mrcnn_mask_conv3    (TimeDistributed)
mrcnn_mask_bn3      (TimeDistributed)
mrcnn_class_conv2   (TimeDistributed)
mrcnn_class_bn2     (TimeDistributed)
mrcnn_mask_conv4    (TimeDistributed)
mrcnn_mask_bn4      (TimeDistributed)
mrcnn_bbox_fc       (TimeDistributed)
mrcnn_mask_deconv   (TimeDistributed)
mrcnn_class_logits  (TimeDistributed)
mrcnn_mask          (TimeDistributed)
```

Epoch 1/1

100/100 [=====] - 105s 1s/step - loss: 0.8017 - rpn_class_loss: 0.0051 - rpn_bbox_loss: 0.3857 - mrcnn_class_loss: 0.0433 - mrcnn_bbox_loss: 0.1242 - mrcnn_mask_loss: 0.2434 - val_loss: 2.1358 - val_rpn_class_loss: 0.0071 - val_rpn_bbox_loss: 1.4229 - val_mrcnn_class_loss: 0.1670 - val_mrcnn_bbox_loss: 0.3696 - val_mrcnn_mask_loss: 0.1693

Epoch 0: Training batch 23 out of 40

Re-starting from epoch 1

Starting at epoch 0. LR=8.862938119652506e-05

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```
fpn_c5p5          (Conv2D)
fpn_c4p4          (Conv2D)
fpn_c3p3          (Conv2D)
fpn_c2p2          (Conv2D)
fpn_p5            (Conv2D)
fpn_p2            (Conv2D)
fpn_p3            (Conv2D)
fpn_p4            (Conv2D)
```

```

In model: rpn_model
    rpn_conv_shared      (Conv2D)
    rpn_class_raw        (Conv2D)
    rpn_bbox_pred        (Conv2D)
mrcnn_mask_conv1        (TimeDistributed)
mrcnn_mask_bn1          (TimeDistributed)
mrcnn_mask_conv2        (TimeDistributed)
mrcnn_mask_bn2          (TimeDistributed)
mrcnn_class_conv1       (TimeDistributed)
mrcnn_class_bn1         (TimeDistributed)
mrcnn_mask_conv3        (TimeDistributed)
mrcnn_mask_bn3          (TimeDistributed)
mrcnn_class_conv2       (TimeDistributed)
mrcnn_class_bn2         (TimeDistributed)
mrcnn_mask_conv4        (TimeDistributed)
mrcnn_mask_bn4          (TimeDistributed)
mrcnn_bbox_fc           (TimeDistributed)
mrcnn_mask_deconv       (TimeDistributed)
mrcnn_class_logits      (TimeDistributed)
mrcnn_mask              (TimeDistributed)

```

Epoch 1/1

100/100 [=====] - 104s 1s/step - loss: 1.188
 9 - rpn_class_loss: 0.0048 - rpn_bbox_loss: 0.7858 - mrcnn_class_loss: 0.0406 - mrcnn_bbox_loss: 0.1369 - mrcnn_mask_loss: 0.2208 - val_loss: 1.2412 - val_rpn_class_loss: 9.9534e-04 - val_rpn_bbox_loss: 0.8403 - val_mrcnn_class_loss: 0.0347 - val_mrcnn_bbox_loss: 0.1667 - val_mrcnn_mask_loss: 0.1985

Epoch 0: Training batch 24 out of 40

Re-starting from epoch 1

Starting at epoch 0. LR=7.976644307687256e-05

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```

fpn_c5p5      (Conv2D)
fpn_c4p4      (Conv2D)
fpn_c3p3      (Conv2D)
fpn_c2p2      (Conv2D)
fpn_p5        (Conv2D)
fpn_p2        (Conv2D)
fpn_p3        (Conv2D)
fpn_p4        (Conv2D)

```

```

In model: rpn_model
    rpn_conv_shared      (Conv2D)
    rpn_class_raw        (Conv2D)
    rpn_bbox_pred        (Conv2D)
mrcnn_mask_conv1        (TimeDistributed)
mrcnn_mask_bn1          (TimeDistributed)
mrcnn_mask_conv2        (TimeDistributed)
mrcnn_mask_bn2          (TimeDistributed)
mrcnn_class_conv1       (TimeDistributed)
mrcnn_class_bn1         (TimeDistributed)
mrcnn_mask_conv3        (TimeDistributed)
mrcnn_mask_bn3          (TimeDistributed)
mrcnn_class_conv2       (TimeDistributed)
mrcnn_class_bn2         (TimeDistributed)

```

```

mrcnn_mask_conv4      (TimeDistributed)
mrcnn_mask_bn4        (TimeDistributed)
mrcnn_bbox_fc         (TimeDistributed)
mrcnn_mask_deconv     (TimeDistributed)
mrcnn_class_logits    (TimeDistributed)
mrcnn_mask            (TimeDistributed)
Epoch 1/1
100/100 [=====] - 104s 1s/step - loss: 1.148
8 - rpn_class_loss: 0.0055 - rpn_bbox_loss: 0.7121 - mrcnn_class_loss: 0.0488 - mrcnn_bbox_loss: 0.1810 - mrcnn_mask_loss: 0.2014 - val_loss: 0.9611 - val_rpn_class_loss: 0.0021 - val_rpn_bbox_loss: 0.4818 - val_mrcnn_class_loss: 0.0428 - val_mrcnn_bbox_loss: 0.2446 - val_mrcnn_mask_loss: 0.1898
Epoch 0: Training batch 25 out of 40
Re-starting from epoch 1

```

Starting at epoch 0. LR=7.17897987691853e-05

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

```

fpn_c5p5      (Conv2D)
fpn_c4p4      (Conv2D)
fpn_c3p3      (Conv2D)
fpn_c2p2      (Conv2D)
fpn_p5        (Conv2D)
fpn_p2        (Conv2D)
fpn_p3        (Conv2D)
fpn_p4        (Conv2D)

```

In model: rpn_model

```

    rpn_conv_shared      (Conv2D)
    rpn_class_raw        (Conv2D)
    rpn_bbox_pred        (Conv2D)
mrcnn_mask_conv1      (TimeDistributed)
mrcnn_mask_bn1        (TimeDistributed)
mrcnn_mask_conv2      (TimeDistributed)
mrcnn_mask_bn2        (TimeDistributed)
mrcnn_class_conv1     (TimeDistributed)
mrcnn_class_bn1       (TimeDistributed)
mrcnn_mask_conv3      (TimeDistributed)
mrcnn_mask_bn3        (TimeDistributed)
mrcnn_class_conv2     (TimeDistributed)
mrcnn_class_bn2       (TimeDistributed)
mrcnn_mask_conv4      (TimeDistributed)
mrcnn_mask_bn4        (TimeDistributed)
mrcnn_bbox_fc         (TimeDistributed)
mrcnn_mask_deconv     (TimeDistributed)
mrcnn_class_logits    (TimeDistributed)
mrcnn_mask            (TimeDistributed)

```

```

Epoch 1/1
100/100 [=====] - 109s 1s/step - loss: 1.307
5 - rpn_class_loss: 0.0054 - rpn_bbox_loss: 0.8429 - mrcnn_class_loss: 0.0614 - mrcnn_bbox_loss: 0.2055 - mrcnn_mask_loss: 0.1921 - val_loss: 1.0380 - val_rpn_class_loss: 0.0031 - val_rpn_bbox_loss: 0.4595 - val_mrcnn_class_loss: 0.0794 - val_mrcnn_bbox_loss: 0.2228 - val_mrcnn_mask_loss: 0.2732
Epoch 0: Training batch 26 out of 40

```

Re-starting from epoch 1

Starting at epoch 0. LR=6.461081889226677e-05

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

fpn_c5p5	(Conv2D)
fpn_c4p4	(Conv2D)
fpn_c3p3	(Conv2D)
fpn_c2p2	(Conv2D)
fpn_p5	(Conv2D)
fpn_p2	(Conv2D)
fpn_p3	(Conv2D)
fpn_p4	(Conv2D)

In model: rpn_model

rpn_conv_shared	(Conv2D)
rpn_class_raw	(Conv2D)
rpn_bbox_pred	(Conv2D)
mrcnn_mask_conv1	(TimeDistributed)
mrcnn_mask_bn1	(TimeDistributed)
mrcnn_mask_conv2	(TimeDistributed)
mrcnn_mask_bn2	(TimeDistributed)
mrcnn_class_conv1	(TimeDistributed)
mrcnn_class_bn1	(TimeDistributed)
mrcnn_mask_conv3	(TimeDistributed)
mrcnn_mask_bn3	(TimeDistributed)
mrcnn_class_conv2	(TimeDistributed)
mrcnn_class_bn2	(TimeDistributed)
mrcnn_mask_conv4	(TimeDistributed)
mrcnn_mask_bn4	(TimeDistributed)
mrcnn_bbox_fc	(TimeDistributed)
mrcnn_mask_deconv	(TimeDistributed)
mrcnn_class_logits	(TimeDistributed)
mrcnn_mask	(TimeDistributed)

Epoch 1/1

100/100 [=====] - 112s 1s/step - loss: 0.8694 - rpn_class_loss: 0.0044 - rpn_bbox_loss: 0.4330 - mrcnn_class_loss: 0.0482 - mrcnn_bbox_loss: 0.1438 - mrcnn_mask_loss: 0.2400 - val_loss: 2.0533 - val_rpn_class_loss: 0.0051 - val_rpn_bbox_loss: 1.4719 - val_mrcnn_class_loss: 0.1173 - val_mrcnn_bbox_loss: 0.2459 - val_mrcnn_mask_loss: 0.2132

Epoch 0: Training batch 27 out of 40

Re-starting from epoch 1

Starting at epoch 0. LR=5.81497370030401e-05

Checkpoint Path: /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_{epoch:04d}.h5

Selecting layers to train

fpn_c5p5	(Conv2D)
fpn_c4p4	(Conv2D)
fpn_c3p3	(Conv2D)
fpn_c2p2	(Conv2D)
fpn_p5	(Conv2D)
fpn_p2	(Conv2D)
fpn_p3	(Conv2D)

```

fpn_p4                                (Conv2D)
In model: rpn_model
    rpn_conv_shared                    (Conv2D)
    rpn_class_raw                      (Conv2D)
    rpn_bbox_pred                     (Conv2D)
mrcnn_mask_conv1                      (TimeDistributed)
mrcnn_mask_bn1                       (TimeDistributed)
mrcnn_mask_conv2                      (TimeDistributed)
mrcnn_mask_bn2                       (TimeDistributed)
mrcnn_class_conv1                    (TimeDistributed)
mrcnn_class_bn1                      (TimeDistributed)
mrcnn_mask_conv3                     (TimeDistributed)
mrcnn_mask_bn3                       (TimeDistributed)
mrcnn_class_conv2                    (TimeDistributed)
mrcnn_class_bn2                      (TimeDistributed)
mrcnn_mask_conv4                     (TimeDistributed)
mrcnn_mask_bn4                       (TimeDistributed)
mrcnn_bbox_fc                        (TimeDistributed)
mrcnn_mask_deconv                    (TimeDistributed)
mrcnn_class_logits                   (TimeDistributed)
mrcnn_mask                           (TimeDistributed)
Epoch 1/1
    3/100 [.....] - ETA: 5:37 - loss: 0.5065 -
    rpn_class_loss: 0.0016 - rpn_bbox_loss: 0.2603 - mrcnn_class_loss: 0.
    0039 - mrcnn_bbox_loss: 0.0558 - mrcnn_mask_loss: 0.1849

```

```

-----
-----
ResourceExhaustedError                                Traceback (most recent call
last)
~/anaconda3/lib/python3.6/site-packages/tensorflow/python/client/session.py in _do_call(self, fn, *args)
    1326         try:
-> 1327             return fn(*args)
    1328         except errors.OpError as e:

~/anaconda3/lib/python3.6/site-packages/tensorflow/python/client/session.py in _run_fn(feed_dict, fetch_list, target_list, options, run_metadata)
    1311         return self._call_tf_sessionrun(
-> 1312             options, feed_dict, fetch_list, target_list, run_metadata)
    1313

~/anaconda3/lib/python3.6/site-packages/tensorflow/python/client/session.py in _call_tf_sessionrun(self, options, feed_dict, fetch_list, target_list, run_metadata)
    1419         self._session, options, feed_dict, fetch_list, target_list,
-> 1420             status, run_metadata)
    1421

~/anaconda3/lib/python3.6/site-packages/tensorflow/python/framework/errors_impl.py in __exit__(self, type_arg, value_arg, traceback_arg)
    515         compat.as_text(c_api.TF_Message(self.status.status)),
-> 516         c_api.TF_GetCode(self.status.status))
    517         # Delete the underlying status object from memory otherwise it stays alive

ResourceExhaustedError: OOM when allocating tensor with shape[2,512,160,160] and type float on /job:localhost/replica:0/task:0/device:GPU:0 by allocator GPU_0_bfc
[[Node: training_27/SGD/gradients/rpn_model/rpn_bbox_pred/convolution_grad/Conv2DBackpropInput = Conv2DBackpropInput[T=DT_FLOAT, _class=["loc:@rpn_model/rpn_bbox_pred/convolution"], data_format="NC_HW", dilations=[1, 1, 1, 1], padding="VALID", strides=[1, 1, 1, 1], use_cudnn_on_gpu=true, _device="/job:localhost/replica:0/task:0/device:GPU:0"] (training_27/SGD/gradients/rpn_model/rpn_bbox_pred/convolution_grad/ShapeN, rpn_bbox_pred/kernel/read, training_27/SGD/gradients/rpn_model/rpn_bbox_pred/convolution_grad/Conv2DBackpropInput-2-TransposeNHWCtoNCHW-LayoutOptimizer)]]
Hint: If you want to see a list of allocated tensors when OOM happens, add report_tensor_allocations_upon_oom to RunOptions for current allocation info.

[[Node: training_27/SGD/gradients/mrcnn_mask_conv1/convolution_grad/Conv2DBackpropInput-0-0-TransposeNCHWtoNHWC-LayoutOptimizer/_41253 = _Recv[client_terminated=false, recv_device="/job:localhost/replica:0/task:0/device:CPU:0", send_device="/job:localhost/replica:0/task:0/device:GPU:0", send_device_incarnation=1, tensor_name="edge_7667_...toOptimizer", tensor_type=DT_FLOAT, _device="/job:localhost/replica:0/task:0/device:CPU:0"]()] ]

```


Hint: If you want to see a list of allocated tensors when OOM happens, add `report_tensor_allocations_upon_oom` to `RunOptions` for current allocation info.

During handling of the above exception, another exception occurred:

```
ResourceExhaustedError                                Traceback (most recent call
last)
<ipython-input-10-be59f87163bb> in <module>()
      22             learning_rate=lr,
      23             epochs=1,
--> 24             layers='heads')
      25
      26         lr = lr * 0.90

~/Documents/PedNet/mrcnn/model.py in train(self, train_dataset, val_d
ataset, learning_rate, epochs, layers, augmentation)
     2312             max_queue_size=100,
     2313             workers=workers,
-> 2314             use_multiprocessing=False,
     2315         )
     2316         self.epoch = max(self.epoch, epochs)

~/anaconda3/lib/python3.6/site-packages/keras/legacy/interfaces.py in
wrapper(*args, **kwargs)
      89             warnings.warn('Update your ``' + object_name +
      90                             `` call to the Keras 2 API: ' +
signature, stacklevel=2)
--> 91             return func(*args, **kwargs)
      92         wrapper._original_function = func
      93         return wrapper

~/anaconda3/lib/python3.6/site-packages/keras/engine/training.py in f
it_generator(self, generator, steps_per_epoch, epochs, verbose, callb
acks, validation_data, validation_steps, class_weight, max_queue_siz
e, workers, use_multiprocessing, shuffle, initial_epoch)
     2222             outs = self.train_on_batch(x, y,
     2223                                         sample_weight=
sample_weight,
-> 2224                                         class_weight=c
lass_weight)
     2225
     2226             if not isinstance(outs, list):

~/anaconda3/lib/python3.6/site-packages/keras/engine/training.py in t
rain_on_batch(self, x, y, sample_weight, class_weight)
     1881             ins = x + y + sample_weights
     1882             self._make_train_function()
-> 1883             outputs = self.train_function(ins)
     1884             if len(outputs) == 1:
     1885                 return outputs[0]

~/anaconda3/lib/python3.6/site-packages/keras/backend/tensorflow_back
end.py in __call__(self, inputs)
     2476             session = get_session()
     2477             updated = session.run(fetches=fetches, feed_dict=feed
```

```

_dict,
-> 2478                                     **self.session_kwargs)
    2479         return updated[:len(self.outputs)]
    2480

~/anaconda3/lib/python3.6/site-packages/tensorflow/python/client/session.py in run(self, fetches, feed_dict, options, run_metadata)
    903     try:
    904         result = self._run(None, fetches, feed_dict, options_ptr,
--> 905                                run_metadata_ptr)
    906         if run_metadata:
    907             proto_data = tf_session.TF_GetBuffer(run_metadata_ptr)
)

~/anaconda3/lib/python3.6/site-packages/tensorflow/python/client/session.py in _run(self, handle, fetches, feed_dict, options, run_metadata)
    1138     if final_fetches or final_targets or (handle and feed_dict_tensor):
    1139         results = self._do_run(handle, final_targets, final_fetches,
-> 1140                                feed_dict_tensor, options, run_metadata)
    1141     else:
    1142         results = []

~/anaconda3/lib/python3.6/site-packages/tensorflow/python/client/session.py in _do_run(self, handle, target_list, fetch_list, feed_dict, options, run_metadata)
    1319     if handle is None:
    1320         return self._do_call(_run_fn, feeds, fetches, targets,
options,
-> 1321                                run_metadata)
    1322     else:
    1323         return self._do_call(_prun_fn, handle, feeds, fetches)

~/anaconda3/lib/python3.6/site-packages/tensorflow/python/client/session.py in _do_call(self, fn, *args)
    1338     except KeyError:
    1339         pass
-> 1340     raise type(e)(node_def, op, message)
    1341
    1342     def _extend_graph(self):

```

ResourceExhaustedError: OOM when allocating tensor with shape[2,512,160,160] and type float on /job:localhost/replica:0/task:0/device:GPU:0 by allocator GPU_0_bfc

```

[[Node: training_27/SGD/gradients/rpn_model/rpn_bbox_pred/convolution_grad/Conv2DBackpropInput = Conv2DBackpropInput[T=DT_FLOAT,
_class=["loc:@rpn_model/rpn_bbox_pred/convolution"], data_format="NC
HW", dilations=[1, 1, 1, 1], padding="VALID", strides=[1, 1, 1, 1], use_cudnn_on_gpu=true, _device="/job:localhost/replica:0/task:0/device:GPU:0"](training_27/SGD/gradients/rpn_model/rpn_bbox_pred/convolution_grad/ShapeN, rpn_bbox_pred/kernel/read, training_27/SGD/gradients/rpn_model/rpn_bbox_pred/convolution_grad/Conv2DBackpropInput-2-TransposeNHWCtoNCHW-LayoutOptimizer)]]

```

Hint: If you want to see a list of allocated tensors when OOM happens, add `report_tensor_allocations_upon_oom` to `RunOptions` for current allocation info.

```
[[Node: training_27/SGD/gradients/mrcnn_mask_conv1/convolution_grad/Conv2DBackpropInput-0-0-TransposeNCHWToNHWC-LayoutOptimizer/_41253 = _Recv[client_terminated=false, recv_device="/job:localhost/replica:0/task:0/device:CPU:0", send_device="/job:localhost/replica:0/task:0/device:GPU:0", send_device_incarnation=1, tensor_name="edge_7667_...tOptimizer", tensor_type=DT_FLOAT, _device="/job:localhost/replica:0/task:0/device:CPU:0"]()] ]]
```

Hint: If you want to see a list of allocated tensors when OOM happens, add `report_tensor_allocations_upon_oom` to `RunOptions` for current allocation info.

Caused by op 'training_27/SGD/gradients/rpn_model/rpn_bbox_pred/convolution_grad/Conv2DBackpropInput', defined at:

```
File "/home/paperspace/anaconda3/lib/python3.6/runpy.py", line 193,
in _run_module_as_main
    "__main__", mod_spec)
File "/home/paperspace/anaconda3/lib/python3.6/runpy.py", line 85,
in _run_code
    exec(code, run_globals)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py", line 16, in <module>
    app.launch_new_instance()
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/traitlets/config/application.py", line 658, in launch_instance
    app.start()
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/ipykernel/kernelapp.py", line 486, in start
    self.io_loop.start()
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tornado/platform/asyncio.py", line 112, in start
    self.asyncio_loop.run_forever()
File "/home/paperspace/anaconda3/lib/python3.6/asyncio/base_events.py", line 421, in run_forever
    self._run_once()
File "/home/paperspace/anaconda3/lib/python3.6/asyncio/base_events.py", line 1426, in _run_once
    handle._run()
File "/home/paperspace/anaconda3/lib/python3.6/asyncio/events.py", line 127, in _run
    self._callback(*self._args)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tornado/platform/asyncio.py", line 102, in _handle_events
    handler_func(fileobj, events)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tornado/stack_context.py", line 276, in null_wrapper
    return fn(*args, **kwargs)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/zmq/eventloop/zmqstream.py", line 450, in _handle_events
    self._handle_recv()
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/zmq/eventloop/zmqstream.py", line 480, in _handle_recv
    self._run_callback(callback, msg)
```

```

File "/home/paperspace/anaconda3/lib/python3.6/site-packages/zmq/ev
entloop/zmqstream.py", line 432, in _run_callback
    callback(*args, **kwargs)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tornad
o/stack_context.py", line 276, in null_wrapper
    return fn(*args, **kwargs)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/ipyker
nel/kernelbase.py", line 283, in dispatcher
    return self.dispatch_shell(stream, msg)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/ipyker
nel/kernelbase.py", line 233, in dispatch_shell
    handler(stream, idents, msg)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/ipyker
nel/kernelbase.py", line 399, in execute_request
    user_expressions, allow_stdin)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/ipyker
nel/ipkernel.py", line 208, in do_execute
    res = shell.run_cell(code, store_history=store_history, silent=si
lent)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/ipyker
nel/zmqshell.py", line 537, in run_cell
    return super(ZMQInteractiveShell, self).run_cell(*args, **kwargs)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/IPytha
n/core/interactiveshell.py", line 2662, in run_cell
    raw_cell, store_history, silent, shell_futures)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/IPytha
n/core/interactiveshell.py", line 2785, in _run_cell
    interactivity=interactivity, compiler=compiler, result=result)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/IPytha
n/core/interactiveshell.py", line 2903, in run_ast_nodes
    if self.run_code(code, result):
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/IPytha
n/core/interactiveshell.py", line 2963, in run_code
    exec(code_obj, self.user_global_ns, self.user_ns)
File "<ipython-input-10-be59f87163bb>", line 24, in <module>
    layers='heads')
File "/home/paperspace/Documents/PedNet/mrcnn/model.py", line 2314,
in train
    use_multiprocessing=False,
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/keras/
legacy/interfaces.py", line 91, in wrapper
    return func(*args, **kwargs)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/keras/
engine/training.py", line 2080, in fit_generator
    self._make_train_function()
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/keras/
engine/training.py", line 990, in _make_train_function
    loss=self.total_loss)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/keras/
legacy/interfaces.py", line 91, in wrapper
    return func(*args, **kwargs)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/keras/
optimizers.py", line 173, in get_updates
    grads = self.get_gradients(loss, params)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/keras/
optimizers.py", line 78, in get_gradients
    grads = K.gradients(loss, params)

```

```

File "/home/paperspace/anaconda3/lib/python3.6/site-packages/keras/
backend/tensorflow_backend.py", line 2515, in gradients
    return tf.gradients(loss, variables, colocate_gradients_with_ops=
True)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tensor
flow/python/ops/gradients_impl.py", line 488, in gradients
    gate_gradients, aggregation_method, stop_gradients)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tensor
flow/python/ops/gradients_impl.py", line 625, in _GradientsHelper
    lambda: grad_fn(op, *out_grads))
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tensor
flow/python/ops/gradients_impl.py", line 379, in _MaybeCompile
    return grad_fn() # Exit early
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tensor
flow/python/ops/gradients_impl.py", line 625, in <lambda>
    lambda: grad_fn(op, *out_grads))
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tensor
flow/python/ops/nn_grad.py", line 514, in _Conv2DGrad
    data_format=data_format),
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tensor
flow/python/ops/gen_nn_ops.py", line 1224, in conv2d_backprop_input
    dilations=dilations, name=name)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tensor
flow/python/framework/op_def_library.py", line 787, in _apply_op_help
er
    op_def=op_def)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tensor
flow/python/framework/ops.py", line 3290, in create_op
    op_def=op_def)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tensor
flow/python/framework/ops.py", line 1654, in __init__
    self._traceback = self._graph._extract_stack() # pylint: disable
=protected-access

```

...which was originally created as op 'rpn_model/rpn_bbox_pred/convol
ution', defined at:

```

File "/home/paperspace/anaconda3/lib/python3.6/runpy.py", line 193,
in _run_module_as_main
    "__main__", mod_spec)
[elided 22 identical lines from previous traceback]
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/IPython
n/core/interactiveshell.py", line 2963, in run_code
    exec(code_obj, self.user_global_ns, self.user_ns)
File "<ipython-input-7-7928c4edfc77>", line 3, in <module>
    model_dir=MODEL_DIR)
File "/home/paperspace/Documents/PedNet/mrcnn/model.py", line 1794,
in __init__
    self.keras_model = self.build(mode=mode, config=config)
File "/home/paperspace/Documents/PedNet/mrcnn/model.py", line 1901,
in build
    layer_outputs.append(rpn([p]))
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/keras/
engine/topology.py", line 619, in __call__
    output = self.call(inputs, **kwargs)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/keras/
engine/topology.py", line 2085, in call
    output_tensors, _, _ = self.run_internal_graph(inputs, masks)

```

```

File "/home/paperspace/anaconda3/lib/python3.6/site-packages/keras/
engine/topology.py", line 2236, in run_internal_graph
    output_tensors = _to_list(layer.call(computed_tensor, **kwargs))
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/keras/
layers/convolutional.py", line 168, in call
    dilation_rate=self.dilation_rate)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/keras/
backend/tensorflow_backend.py", line 3335, in conv2d
    data_format=tf_data_format)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tensor
flow/python/ops/nn_ops.py", line 782, in convolution
    return op(input, filter)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tensor
flow/python/ops/nn_ops.py", line 870, in __call__
    return self.conv_op(inp, filter)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tensor
flow/python/ops/nn_ops.py", line 522, in __call__
    return self.call(inp, filter)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tensor
flow/python/ops/nn_ops.py", line 206, in __call__
    name=self.name)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tensor
flow/python/ops/gen_nn_ops.py", line 953, in conv2d
    data_format=data_format, dilations=dilations, name=name)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tensor
flow/python/framework/op_def_library.py", line 787, in _apply_op_help
er
    op_def=op_def)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tensor
flow/python/framework/ops.py", line 3290, in create_op
    op_def=op_def)
File "/home/paperspace/anaconda3/lib/python3.6/site-packages/tensor
flow/python/framework/ops.py", line 1654, in __init__
    self._traceback = self._graph._extract_stack() # pylint: disable
=protected-access

```

ResourceExhaustedError (see above for traceback): OOM when allocating tensor with shape[2,512,160,160] and type float on /job:localhost/replica:0/task:0/device:GPU:0 by allocator GPU_0_bfc

```

[[Node: training_27/SGD/gradients/rpn_model/rpn_bbox_pred/convolu
tion_grad/Conv2DBackpropInput = Conv2DBackpropInput[T=DT_FLOAT,
_class=["loc:@rpn_model/rpn_bbox_pred/convolution"], data_format="NC
HW", dilations=[1, 1, 1, 1], padding="VALID", strides=[1, 1, 1, 1], u
se_cudnn_on_gpu=true, _device="/job:localhost/replica:0/task:0/devic
e:GPU:0"](training_27/SGD/gradients/rpn_model/rpn_bbox_pred/convoluti
on_grad/ShapeN, rpn_bbox_pred/kernel/read, training_27/SGD/gradients/
rpn_model/rpn_bbox_pred/convolution_grad/Conv2DBackpropInput-2-Transp
oseNHWCtoNCHW-LayoutOptimizer)]]

```

Hint: If you want to see a list of allocated tensors when OOM happens, add report_tensor_allocations_upon_oom to RunOptions for current allocation info.

```

[[Node: training_27/SGD/gradients/mrcnn_mask_conv1/convoluti
on_grad/Conv2DBackpropInput-0-0-TransposeNCHWtoNHWC-LayoutOptimizer/_
41253 = Recv[client_terminated=false, recv_device="/job:localhost/re
plica:0/task:0/device:CPU:0", send_device="/job:localhost/replica:0/t
ask:0/device:GPU:0", send_device_incarnation=1, tensor_name="edge_766

```

```
7_...tOptimizer", tensor_type=DT_FLOAT, _device="/job:localhost/repl
ica:0/task:0/device:CPU:0"]())]]
```

Hint: If you want to see a list of allocated tensors when OOM happens, add `report_tensor_allocations_upon_oom` to `RunOptions` for current allocation info.

```
In [ ]: # Save weights
        # Typically not needed because callbacks save after every epoch
        # Uncomment to save manually
        # model_path = os.path.join(MODEL_DIR, "mask_rcnn_shapes.h5")
        # model.keras_model.save_weights(model_path)
```

Detection

```
In [12]: class InferenceConfig(PedConfig):
        GPU_COUNT = 1
        IMAGES_PER_GPU = 1

        inference_config = InferenceConfig()

        # Recreate the model in inference mode
        model = modellib.MaskRCNN(mode="inference",
                                   config=inference_config,
                                   model_dir=MODEL_DIR)

        # Get path to saved weights
        # Either set a specific path or find last trained weights
        # model_path = os.path.join(ROOT_DIR, ".h5 file name here")
        # model_path = model.find_last()

        # Load trained weights
        print("Loading weights from ", "/home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_0001.h5")
        model.load_weights("/home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_0001.h5", by_name=True)
```


```
Loading weights from /home/paperspace/Documents/PedNet/logs/ped20181117T1835/mask_rcnn_ped_0001.h5
Re-starting from epoch 1
```

```
In [13]: # Test on a random image
# image_id = random.choice(dataset_val.image_ids)
image_id = 2
original_image, image_meta, gt_class_id, gt_bbox, gt_mask =\
    modellib.load_image_gt(dataset_val, inference_config,
                           image_id, use_mini_mask=False)

log("original_image", original_image)
log("image_meta", image_meta)
log("gt_class_id", gt_class_id)
log("gt_bbox", gt_bbox)
log("gt_mask", gt_mask)

visualize.display_instances(original_image, gt_bbox, gt_mask, gt_class_id,
                           dataset_train.class_names, figsize=(8, 8))
```

```
original_image      shape: (640, 640, 3)      min: 0.00000
    max: 255.00000 float64
image_meta          shape: (16,)          min: 0.00000
    max: 640.00000 int64
gt_class_id         shape: (1, 1)         min: 2.00000
    max: 2.00000 int32
gt_bbox            shape: (1, 4)          min: 183.00000
    max: 640.00000 int32
gt_mask            shape: (640, 640, 1)   min: 0.00000
    max: 1.00000 float64
```





```
In [14]: results = model.detect([original_image], verbose=1)

r = results[0]
visualize.display_instances(original_image, r['rois'], r['masks'], r[
    'class_ids'],
                           dataset_val.class_names, r['scores'], ax=
    get_ax())
```

Processing 1 images

image	shape: (640, 640, 3)	min: 0.00000
max: 255.00000 float64		
molded_images	shape: (1, 640, 640, 3)	min: -123.70000
max: 151.10000 float64		
image metas	shape: (1, 16)	min: 0.00000
max: 640.00000 int64		
anchors	shape: (1, 102300, 4)	min: -0.14164
max: 1.04149 float32		



Evaluation

```
In [16]: # Compute VOC-Style mAP @ IoU=0.5
# Running on 10 images. Increase for better accuracy.
image_ids = np.random.choice(dataset_val.image_ids, dataset_val.N)
APs = []
for image_id in image_ids:
    # Load image and ground truth data
    image, image_meta, gt_class_id, gt_bbox, gt_mask = \
        modellib.load_image_gt(dataset_val, inference_config,
                                image_id, use_mini_mask=False)
    molded_images = np.expand_dims(modellib.mold_image(image, inference_config), 0)
    # Run object detection
    results = model.detect([image], verbose=0)
    r = results[0]
    # Compute AP
    AP, precisions, recalls, overlaps = \
        utils.compute_ap(gt_bbox, gt_class_id, gt_mask,
                          r["rois"], r["class_ids"], r["scores"], r['masks'])
    APs.append(AP)

print("mAP: ", np.mean(APs))
```

mAP: 0.9833333333333333

Data Visualization

```
In [5]: config = PedConfig()

if config.NAME == 'ped':
    dataset = PedDataset()
    dataset.load_ped(istrain=True, st_index=0)

# Must call before using the dataset
dataset.prepare()

print("Image Count: {}".format(len(dataset.image_ids)))
print("Class Count: {}".format(dataset.num_classes))
for i, info in enumerate(dataset.class_info):
    print("{:3}. {:50}".format(i, info['name']))
```

Image Count: 20

Class Count: 4

- 0. BG
- 1. bike
- 2. car
- 3. person

```
In [6]: # Load and display random samples
image_ids = np.random.choice(dataset.image_ids, 4)
for image_id in image_ids:
    image = dataset.load_image(image_id)
    mask, class_ids = dataset.load_mask(image_id)
    visualize.display_top_masks(image, mask, class_ids, dataset.class_names)
```



```
In [7]: # Load random image and mask.
image_id = random.choice(dataset.image_ids)
image = dataset.load_image(image_id)
mask, class_ids = dataset.load_mask(image_id)
# Compute Bounding box
bbox = utils.extract_bboxes(mask)

# Display image and additional stats
print("image_id ", image_id, dataset.image_reference(image_id))
log("image", image)
log("mask", mask)
log("class_ids", class_ids)
log("bbox", bbox)
# Display image and instances
visualize.display_instances(image, bbox, mask, class_ids, dataset.class_names)
```

```
image_id  0 (480, 640)
image                                shape: (480, 640, 3)           min:  33.00000
  max: 255.00000 float64
mask                                shape: (480, 640, 1)           min:  0.00000
  max:  1.00000 float64
class_ids                             shape: (1, 1)              min:  1.00000
  max:  1.00000 int32
bbox                                shape: (1, 4)                min:  0.00000
  max: 640.00000 int32
```



```
In [8]: # Load random image and mask.
image_id = np.random.choice(dataset.image_ids, 1)[0]
image = dataset.load_image(image_id)
mask, class_ids = dataset.load_mask(image_id)
original_shape = image.shape
# Resize
image, window, scale, padding, _ = utils.resize_image(
    image,
    min_dim=config.IMAGE_MIN_DIM,
    max_dim=config.IMAGE_MAX_DIM,
    mode=config.IMAGE_RESIZE_MODE)
mask = utils.resize_mask(mask, scale, padding)
# Compute Bounding box
bbox = utils.extract_bboxes(mask)

# Display image and additional stats
print("image_id: ", image_id, dataset.image_reference(image_id))
print("Original shape: ", original_shape)
log("image", image)
log("mask", mask)
log("class_ids", class_ids)
log("bbox", bbox)
# Display image and instances
visualize.display_instances(image, bbox, mask, class_ids, dataset.class_names)
```

```
image_id: 11 (480, 640)
Original shape: (480, 640, 3)
image shape: (640, 640, 3) min: 0.00000
max: 244.00000 float64
mask shape: (640, 640, 1) min: 0.00000
max: 1.00000 float64
class_ids shape: (1, 1) min: 2.00000
max: 2.00000 int32
bbox shape: (1, 4) min: 90.00000
max: 640.00000 int32
```



```

In [9]: # Generate Anchors
backbone_shapes = modellib.compute_backbone_shapes(config, config.IMA
GE_SHAPE)
anchors = utils.generate_pyramid_anchors(config.RPN_ANCHOR_SCALES,
                                          config.RPN_ANCHOR_RATIOS,
                                          backbone_shapes,
                                          config.BACKBONE_STRIDES,
                                          config.RPN_ANCHOR_STRIDE)

# Print summary of anchors
num_levels = len(backbone_shapes)
anchors_per_cell = len(config.RPN_ANCHOR_RATIOS)
print("Count: ", anchors.shape[0])
print("Scales: ", config.RPN_ANCHOR_SCALES)
print("ratios: ", config.RPN_ANCHOR_RATIOS)
print("Anchors per Cell: ", anchors_per_cell)
print("Levels: ", num_levels)
anchors_per_level = []
for l in range(num_levels):
    num_cells = backbone_shapes[l][0] * backbone_shapes[l][1]
    anchors_per_level.append(anchors_per_cell * num_cells // config.R
PN_ANCHOR_STRIDE**2)
    print("Anchors in Level {}: {}".format(l, anchors_per_level[l]))

```

```

Count: 102300
Scales: (8, 16, 32, 64, 128)
ratios: [0.5, 1, 2]
Anchors per Cell: 3
Levels: 5
Anchors in Level 0: 76800
Anchors in Level 1: 19200
Anchors in Level 2: 4800
Anchors in Level 3: 1200
Anchors in Level 4: 300

```

```

In [10]: import matplotlib.patches as patches
import matplotlib.lines as lines
from matplotlib.patches import Polygon

## Visualize anchors of one cell at the center of the feature map of
a specific level


# Load and draw random image
image_id = np.random.choice(dataset.image_ids, 1)[0]
image, image_meta, _, _, _ = modellib.load_image_gt(dataset, config,
image_id)
fig, ax = plt.subplots(1, figsize=(10, 10))
ax.imshow(image.astype(np.uint8))
levels = len(backbone_shapes)

for level in range(levels):
    colors = visualize.random_colors(levels)
    # Compute the index of the anchors at the center of the image
    level_start = sum(anchors_per_level[:level]) # sum of anchors of
previous levels
    level_anchors = anchors[level_start:level_start+anchors_per_level
[level]]
    print("Level {}. Anchors: {:6} Feature map Shape: {}".format(level,
level_anchors.shape[0],
                                                                    backbone_shapes[level]))
    center_cell = backbone_shapes[level] // 2
    center_cell_index = (center_cell[0] * backbone_shapes[level][1] +
center_cell[1])
    level_center = center_cell_index * anchors_per_cell
    center_anchor = anchors_per_cell * (
        (center_cell[0] * backbone_shapes[level][1] / config.RPN_ANCH
OR_STRIDE**2) \
        + center_cell[1] / config.RPN_ANCHOR_STRIDE)
    level_center = int(center_anchor)

    # Draw anchors. Brightness show the order in the array, dark to b
right.
    for i, rect in enumerate(level_anchors[level_center:level_center+
anchors_per_cell]):
        y1, x1, y2, x2 = rect
        p = patches.Rectangle((x1, y1), x2-x1, y2-y1, linewidth=2, fa
cecolor='none',
                                                                    edgecolor=(i+1)*np.array(colors[level])
/ anchors_per_cell)
        ax.add_patch(p)

```


Level 0. Anchors: 76800 Feature map Shape: [160 160]
 Level 1. Anchors: 19200 Feature map Shape: [80 80]
 Level 2. Anchors: 4800 Feature map Shape: [40 40]
 Level 3. Anchors: 1200 Feature map Shape: [20 20]
 Level 4. Anchors: 300 Feature map Shape: [10 10]



Create data generator

```
In [11]: random_rois = 2000
g = modellib.DataGenerator(
    dataset, config, shuffle=True, random_rois=random_rois,
    batch_size=4,
    detection_targets=True)
g
```

```
Out[11]: <mrcnn.model.DataGenerator at 0x7f973101bc88>
```

```

In [13]: # Get Next Image
if random_rois:
    [normalized_images, image_meta, rpn_match, rpn_bbox, gt_class_ids
    , gt_boxes, gt_masks, rpn_rois, rois], \
    [mrcnn_class_ids, mrcnn_bbox, mrcnn_mask] = g.__getitem__(0)

    log("rois", rois)
    log("mrcnn_class_ids", mrcnn_class_ids)
    log("mrcnn_bbox", mrcnn_bbox)
    log("mrcnn_mask", mrcnn_mask)
else:
    [normalized_images, image_meta, rpn_match, rpn_bbox, gt_boxes, gt
    _masks], _ = next(g)

    log("gt_class_ids", gt_class_ids)
    log("gt_boxes", gt_boxes)
    log("gt_masks", gt_masks)
    log("rpn_match", rpn_match, )
    log("rpn_bbox", rpn_bbox)
    image_id = modellib.parse_image_meta(image_meta)["image_id"][0]
    print("image_id: ", image_id, dataset.image_reference(image_id))

# Remove the last dim in mrcnn_class_ids. It's only added
# to satisfy Keras restriction on target shape.
mrcnn_class_ids = mrcnn_class_ids[:, :, 0]

```

```

rois                shape: (4, 32, 4)           min:    7.00000
  max: 639.00000  int32
mrcnn_class_ids     shape: (4, 32, 1, 1)       min:    0.00000
  max:   2.00000  int32
mrcnn_bbox          shape: (4, 32, 4, 4)       min:   -3.31551
  max:   2.68606  float32
mrcnn_mask          shape: (4, 32, 28, 28, 4)  min:    0.00000
  max:   1.00000  float32
gt_class_ids        shape: (4, 100)           min:    0.00000
  max:   2.00000  int32
gt_boxes            shape: (4, 100, 4)        min:    0.00000
  max: 640.00000  int32
gt_masks            shape: (4, 640, 640, 100)  min:    0.00000
  max:   1.00000  bool
rpn_match           shape: (4, 102300, 1)     min:   -1.00000
  max:   1.00000  int32
rpn_bbox            shape: (4, 256, 4)        min:   -7.07107
  max:   7.41752  float64
image_id: 3 (480, 640)

```

```
In [19]: b = 0

# Restore original image (reverse normalization)
sample_image = modellib.unmold_image(normalized_images[b], config)

# Compute anchor shifts.
indices = np.where(rpn_match[b] == 1)[0]
refined_anchors = utils.apply_box_deltas(anchors[indices], rpn_bbox[b
, :len(indices)] * config.RPN_BBOX_STD_DEV)
log("anchors", anchors)
log("refined_anchors", refined_anchors)

# Get list of positive anchors
positive_anchor_ids = np.where(rpn_match[b] == 1)[0]
print("Positive anchors: {}".format(len(positive_anchor_ids)))
negative_anchor_ids = np.where(rpn_match[b] == -1)[0]
print("Negative anchors: {}".format(len(negative_anchor_ids)))
neutral_anchor_ids = np.where(rpn_match[b] == 0)[0]
print("Neutral anchors: {}".format(len(neutral_anchor_ids)))


# ROI breakdown by class
for c, n in zip(dataset.class_names, np.bincount(mrcnn_class_ids[b].f
latten())):
    if n:
        print("{:23}: {}".format(c[:20], n))

# Show positive anchors
visualize.draw_boxes(sample_image, boxes=anchors[positive_anchor_ids
],
                    refined_boxes=refined_anchors)
```

```
anchors          shape: (102300, 4)          min: -90.50967
  max: 666.50967  float64
refined_anchors  shape: (1, 4)              min: 364.00000
  max: 437.00000  float32
Positive anchors: 1
Negative anchors: 255
Neutral anchors: 102044
BG                : 22
CAR               : 10
```



```
In [20]: visualize.draw_boxes(sample_image, boxes=anchors[negative_anchor_ids  
])
```



```
In [21]: visualize.draw_boxes(sample_image, boxes=anchors[np.random.choice(neu  
tral_anchor_ids, 100)])
```

